

STARTER TEMPLATE · ADAPT BEFORE ADOPTION

The Bank AI Workflow SOP Template

A fillable operating procedure for documenting AI-assisted work, human review, data handling, vendor controls, monitoring, and shutoff triggers.

WHO IT IS FOR

Operations, compliance, lending, any team running AI-assisted work

FILL TIME

18 min

WHAT'S INSIDE

- 01 Part 1: blank AI workflow SOP template
- 02 Workflow identity, scope, and prohibited uses
- 03 Human review, decision controls, and model-risk considerations
- 04 Vendor and third-party lifecycle
- 05 Monitoring thresholds and shutoff triggers
- 06 Governance, incidents, recovery, and retirement
- 07 Part 2: completed example - commercial loan memo drafting
- 08 What good looks like and next step

SECTION 01

Part 1: blank AI workflow SOP template

Complete this page first. It turns the SOP from guidance into a working control record that risk, compliance, audit, and exam teams can inspect.

Fillable front page

Field	Response
Workflow ID	[AI-__-__]
Workflow Name	[Plain-English workflow name]
Department	[Business line or function]
Named Owner	[Person accountable for operation]
Accountable Executive	[Senior executive accountable for risk]
Tool/Vendor	[Approved tool, vendor, version if known]
Deployment Type	[Public AI / enterprise account / vendor feature / private deployment / internal model]
Data Class	[Public / internal / confidential-NPI / regulated / exam-sensitive / privileged]
Customer Data Involved	[No / yes / de-identified / approved private environment only]
Regulated Process Involved	[Credit / BSA-AML / fraud / complaints / regulatory reporting / none]
Risk Tier	[Low / moderate / high / blocked]
Approval Date	[YYYY-MM-DD]
Next Review Date	[YYYY-MM-DD]
Material Error Threshold	[X% or other defined limit]
Customer-Impacting Error Threshold	[Any occurrence / X occurrences]
Shutoff Threshold	[Severe event / vendor model change / two consecutive months over threshold]
Fallback Process	[Manual process or alternate approved workflow]
Records Retained	[Prompt, output, review note, approval, ticket, vendor evidence, or final artifact]
Last Version Reviewed	[Version/date/reviewer]

SECTION 02

Workflow identity, scope, and prohibited uses

Register one workflow, not a whole tool. The SOP should let a reader state what the workflow does, what it cannot do, who reviews it, what evidence is retained, and when use is paused.

- Business purpose and output - name the job and the artifact. Example: "Generate a first-draft procedure summary for human review; output is a draft staff memo, not an official record until approved."
- In scope - list the exact tasks the AI may perform, such as drafting, summarizing, comparing, extracting, or formatting approved source material.
- Out of scope / prohibited - adaptable clause language: "This workflow does not make or recommend credit decisions, assign risk ratings, determine customer eligibility, calculate official figures, communicate directly with customers, or serve as a system of record unless separately approved."
- Data boundary - list allowed inputs and prohibited inputs. Customer NPI, account data, exam-sensitive material, SAR/AML information, privileged material, security-control data, credentials, and full identifiers require explicit approval and an approved environment.
- Evidence retained - identify the retained record before go-live: final approved output, reviewer note, approval ticket, exception log entry, vendor review, monitoring report, or other evidence.

SECTION 03

Human review, decision controls, and model-risk considerations

Where an AI-assisted workflow supports quantitative, customer-impacting, or regulated decision processes, the institution applies model-risk, compliance, validation, monitoring, and effective-challenge controls appropriate to the use case. Generative AI workflows used for drafting, summarization, or workflow support should be governed through approved-tool controls, data restrictions, human review, documentation, vendor oversight, monitoring, and incident-response procedures.

- Human-in-the-loop control - adaptable clause language: "A qualified reviewer checks every AI output for factual accuracy, missing context, unsupported statements, stale information, data leakage, and tone before the output is used or retained."
- Decision-control boundary - state whether the output can affect credit availability, pricing, fraud handling, BSA/AML work, complaints, regulatory reporting, or any other regulated process. If yes, require Compliance, Risk, and business-owner review before production use.
- Effective challenge - name the reviewer who can reject the output, override it, require manual work, or pause the workflow without business-line retaliation.
- Credit/adverse-action control - AI-generated adverse-action language may not be used unless Compliance confirms that the stated reasons are specific, accurate, and traceable to factors actually considered.
- OCC scope note - do not treat every generative AI drafting workflow as a formal model-validation event. Apply current model-risk principles where the workflow informs quantitative, customer-impacting, or regulated decisions, and apply documented governance controls for drafting and workflow-support use cases.

SECTION 04

Vendor and third-party lifecycle

If the workflow uses a vendor, embedded AI feature, hosted model, or enterprise AI platform, document the third-party lifecycle. Third-party use does not move responsibility away from the bank.

AI vendor lifecycle controls

Lifecycle stage	AI-specific question	Evidence retained
Planning	Why is AI needed for this workflow, and what manual process remains available?	Business case, risk tier, fallback process
Due diligence	Does the vendor train on bank or customer data, and what security, model, and data-use evidence was reviewed?	SOC report, security review, AI/data-use documentation, financial review
Contracting	Are data use, deletion, breach notice, subcontractors, audit rights, regulator access, and retention addressed?	Contract terms, DPA, legal/compliance approval
Monitoring	Who reviews tool changes, model/version updates, errors, incidents, and performance?	Monitoring report, change log, incident log, vendor review
Termination	How is access disabled and bank/customer data returned or deleted?	Termination checklist, access removal, deletion/return confirmation

SECTION 05

Monitoring thresholds and shutoff triggers

Do not write "monitor as needed." Put the actual pause thresholds in the SOP so the owner knows when use must stop and who can restart it.

- Pre-deployment check - document the initial test that established fitness for the approved scope. For higher-risk workflows, require independent validation or model-risk review before production use.
- Ongoing metrics - choose 2-4 measures such as factual-error rate, material rework rate, data-handling exceptions, customer-impacting errors, incident count, and vendor change events.
- Escalation - adaptable clause language: "If a pause trigger occurs, the workflow owner stops use, notifies the accountable executive and control partners, preserves evidence, and documents the restart decision before use resumes."

Required monitoring fields

Control field	Required entry
Material error threshold	[X%] of reviewed outputs or [X] material errors per month
Customer-impacting error threshold	[Any occurrence] or [X occurrences]
Data-handling exception threshold	[Any restricted data entered into an unapproved tool]
Pause trigger	[Two consecutive months over threshold / any severe event / vendor model change / data exposure]
Restart approval	[Owner + Compliance + Risk / committee / accountable executive]
Monitoring cadence	[Monthly / quarterly / event-triggered]

SECTION 06

Governance, incidents, recovery, and retirement

Tie the SOP to existing policies and decide in advance what happens when the AI gets it wrong, changes materially, or no longer belongs in production.

- Approvals on record - list approval date, approver, policy linkage, risk acceptance, and any conditions before production use.
- Change management - adaptable clause language: "Material changes to inputs, scope, output, model/version, vendor, data class, or regulated-process impact require re-review before continued use."
- Incident response - define same-business-day notification for material errors, customer-impacting outputs, data exposure, prohibited inputs, suspicious activity concerns, or examiner-facing errors.
- Recovery / fallback - confirm the manual path, reprocess affected items through human review, and preserve evidence of correction.
- Retirement - disable access, confirm data return/deletion where applicable, archive the SOP, mark the use case retired, and record the reason and date.

SECTION 07

Part 2: completed example - commercial loan memo drafting

Use this as an example only. It separates the lending scenario from the blank template and corrects the credit-control language that belongs in any lending workflow.

- Workflow ID/name - AI-LEND-001: AI-assisted commercial loan memo drafting.
- Purpose/output - generate first-draft narrative sections of a commercial credit memo from approved underwriting inputs; the output is a draft for analyst editing and credit officer approval.
- Credit boundary - commercial credit workflow; ECOA / Regulation B and fair-lending considerations still apply. The AI does not recommend approval or denial, assign risk ratings, draft adverse-action reasons, or determine credit terms. Any use that could affect credit availability, terms, or reasons for action requires Compliance review.
- Data boundary - borrower financials, spreading output, prior approved memos, and analyst-selected public industry data may be used only inside the approved vendor environment. Full SSNs, account numbers, consumer credit-report data, SAR/AML information, and unapproved customer NPI are prohibited inputs.
- Human review - analyst reviews every generated section line by line; credit officer signs off before the memo advances; Compliance reviews any change that could affect adverse-action, fair-lending, or customer-facing explanations.
- Thresholds - material error threshold [X%]; customer-impacting error threshold [any occurrence]; pause trigger [vendor model change, data exposure, or two consecutive months above threshold]; restart approval [owner + Compliance + Risk].

SECTION 08

What good looks like and next step

A credible SOP is not a policy essay. It is a control record that shows the workflow is bounded, reviewed, monitored, and shut off when needed.

- What good looks like - a reader can tell what the AI does, what it cannot do, who reviews every output, which data is allowed, which evidence is retained, and what shuts it off.
- Common mistake - using "approved for AI" language without a specific workflow, owner, data class, or pause threshold.
- Common mistake - copying clause language without tailoring it to the bank's tool stack, retention schedule, approval workflow, and risk appetite.
- Next step: run a 45-minute AI Workflow SOP Build Session for one live workflow. Leave with the filled front page, data boundary, reviewer standard, vendor evidence list, monitoring thresholds, and first inventory row.

SOURCE BASIS

- SR 26-2 and OCC Bulletin 2026-13 - revised model-risk guidance, scope limitations for generative and agentic AI, and risk-based principles where applicable
- NIST AI Risk Management Framework (AI RMF 1.0): Govern, Map, Measure, Manage
- Interagency Guidance on Third-Party Relationships: Risk Management (FRB / FDIC / OCC, June 2023)
- ECOA / Regulation B and CFPB Circular 2023-03 for credit and adverse-action controls
- AIEOG AI Lexicon and Financial Services AI Risk Management Framework - US Treasury / FBIIC / FSSCC (Feb 2026)

ADAPT BEFORE ADOPTION

This is a starter template, not legal advice or an examiner-approved policy. Replace every [bracketed] placeholder, align the language with your institution's policies, and route the final document through compliance, risk, and legal review.