SKILLS MASTER CLASS

AIDB Operators Cut — Companion Experience

Nufar Gaspar × Nathaniel Whittemore

SKILL POWER 0%

APPRENTICE BUILDER ARSENAL STRATEGIST ARCHITECT

1 APPRENTICE

Understanding the Skill Landscape

WHAT ARE SKILLS?

At their core, skills are **folders** — not just markdown files — containing instructions, scripts, and resources that give AI agents new capabilities.

- Skills can **run code**, **trigger external systems**, **call APIs**, and **orchestrate multi-step workflows**
- **Two modes:** agents discover & invoke them autonomously, OR humans trigger them manually (/research this topic). Skills serve both.
- The problem they solve: system prompts become bloated → performance degrades. Skills load the right knowledge dynamically, only when needed.

THE PORTABILITY BREAKTHROUGH

Custom GPTs were locked inside ChatGPT. Skills are **portable, human-readable markdown folders**. Build once, use across tools. No proprietary format, no vendor lock-in. Anyone can read, edit, and maintain a skill — no engineering required.

This is what custom GPTs *should* have been.

THE ECOSYSTEM

44+ AI tools already support skills: Claude Code, Cursor, Windsurf, GitHub Copilot, Codex, Gemini CLI, Notion, and many more. This is becoming a universal standard.

TWO FUNDAMENTAL TYPES

TYPE	WHAT IT DOES	DURABILITY
CAPABILITY UPLIFT	Enables new functions the model can't do well on its own	May become obsolete as models improve
ENCODED PREFERENCE	Sequences existing capabilities according to YOUR workflow	Gets more valuable over time

Spend most time on **preference skills**. They encode how YOUR team works.

PROGRESSIVE DISCLOSURE

LAYER	WHAT LOADS	TOKEN COST
1. Description	~100 tokens in system prompt	Always loaded
2. SKILL.md body	Full instructions	Only when triggered
3. Folder contents	Scripts, assets, references	Only when needed

> ⚠ Security Warning
> Third-party skills are code that runs with your agent's permissions. Always audit before installing. Treat skill installation like installing a browser extension.

BUILDER

WHEN TO BUILD A SKILL

- You do something **more than 3 times** with an AI tool
- You keep pasting the same instructions repeatedly
- You want **consistent, reliable output** across sessions
- You want to **enforce standardization** — the skill becomes the single source of truth
- **Growth mindset:** Build for things you *couldn't* do before, not just automating what you already do

> **SCOPING RULE**
> One clear job per skill. If you can't describe what it does in one sentence, it's probably two skills.

REUSE VS. CREATE

Growing libraries exist (Anthropic skills repo: 87K+ GitHub stars). But for your **first skills: build your own.** Browsing marketplaces is tedious, quality varies, and security concerns are real. Building from scratch teaches you the craft faster than adopting someone else's 70%-right skill.

THE SKILL CREATOR TOOL

Claude's built-in tool — but the pattern can be emulated in any tool by building a "skill for building skills." It interviews you, runs evals, does A/B testing, and iteratively improves your skill's description for better triggering.

SKILL ANATOMY

name	Lowercase, hyphens, max 64 chars. Gerund form: analyzing-data, preparing-meetings
description	**The most critical line.** A **trigger, not a summary.** Write for the model asking "when should I fire?" Pro tip: make it LOUDER rather than quieter. Write in third person. Include both what it does and when to use it.
Instructions	Favor **numbered steps or bulleted lists over prose.** Set degrees of freedom: tight for fragile ops, loose for creative tasks.

Output Format	Show, don't describe. Include a literal template or example.
Gotcha Section	**Highest-signal content.** Where does the model go wrong? "I know you'll want to do X — don't."
Constraints	What NOT to do. Sharp and specific to THIS skill.

> Skip This
> **Identity / Role section** ("Act as a senior analyst...") — legacy prompt engineering pattern. Tell the model what YOUR approach does differently, not what persona to adopt.

THE 5 SKILL KILLERS

1. **Description doesn't trigger properly** — too vague, too narrow, or wrong person
   Fix: Specific, loud, third-person. "Use when..." format.
2. **Over-defining the process** — railroading instead of guiding
   Fix: Set degrees of freedom. Tight for fragile ops, loose for creative.
3. **Stating the obvious** — wasting tokens on what the model knows
   Fix: Challenge every paragraph: "Does Claude really need this?"
4. **Missing gotcha section** — not capturing failure patterns
   Fix: Document every failure you've seen. This IS the skill's value.
5. **Monolithic blob** — everything in one file
   Fix: SKILL.md under 500 lines. Move references to separate files.

CONTEXT BINDING: IN-FOLDER VS. EXTERNAL

PUT IN THE FOLDER WHEN...	POINT EXTERNALLY WHEN...
Context is specific to this skill and travels with it (rubric, template, persona, examples)	Context is shared across skills or changes independently (CLAUDE.md, project docs, stakeholder list)

Rule of thumb: "about the skill" = inside. "About you/your org" = outside.

SHOWCASE: MEETING PREP SKILL

```
meeting-prep/
├── SKILL.md              # Core instructions
├── stakeholder-context.md  # External pointer
├── output-template.md      # Structured brief format
├── scenarios.md            # Meetings that go off the rails
├── examples.md             # Great vs. mediocre prep
└── skills/
    └── meeting-sim/
        └── SKILL.md        # Nested: simulate the meeting
```

📄 meeting-prep/SKILL.md SKILL.md ▼

📦 Full skill folder (SKILL.md + scenarios.md + output-template.md + meeting-sim/) available in the **Expansion Pack → Blueprints** section below.

3 ARSENAL

Your Starter Kit

▲

4 SKILLS EVERY KNOWLEDGE WORKER SHOULD HAVE

IMPORTANT

These are **starting points, not finished products.** Every skill below needs to be customized to YOUR context: your sources, your stakeholders, your quality bar, your terminology. Copy the structure, then make it yours. The trigger phrases, the gotcha sections, the output formats — all should reflect how YOU actually work.

🔍 research-with-confidence/SKILL.md CAPABILITY ▲

```
---
name: researching-with-confidence
```

```
description: >
  Use when the user says "research [topic]", "what's the latest
  on [topic]", "deep dive into [topic]", "find out about",
  "investigate", or asks any question that requires gathering
  information from multiple sources. Also triggers on "is it
  true that...", "verify this claim", or "fact-check". Runs
  parallel multi-angle research with confidence scoring and
  cross-source fact-checking.
---

# Research with Confidence

## Inputs (confirm with user)

- Topic: [from user]
- Time horizon: last week / last quarter / last year / all time
- Source types:
  Tier 1: McKinsey, HBR, Gartner, peer-reviewed journals
  Tier 2: TechCrunch, industry blogs, company announcements
  Tier 3: Reddit, X/Twitter threads, Hacker News, forums
  Primary: SEC filings, patent databases, government data
```

😈 devils-advocate/SKILL.md PREFERENCE ▲

```
---
name: devils-advocate
description: >
  Use when the user says "devil's advocate this", "poke holes
  in this", "stress test", "what am I missing", "challenge
  this", "what could go wrong", or asks for critical review
  of any proposal, plan, decision, or document. Also triggers
  when user says "before I send this" or "is this solid".
  Systematically identifies hidden assumptions, blind spots,
  biases in both the presenter and the model's own reasoning.
---

# Devil's Advocate

## Steps
```

```
1. Read the target document, decision, or proposal completely
2. Identify 3-5 hidden assumptions the author is making
3. For each assumption, construct the strongest
   counter-argument (steelman the opposition)
4. Find the single biggest blind spot — what's
   not being considered at all?
5. Check for presenter's biases: anchoring,
```

☀️ morning-briefing/SKILL.md PREFERENCE ▲

```
---
name: morning-briefing
description: >
  Use when the user says "morning brief", "start my day",
  "daily briefing", "what's on today", "catch me up", or at
  the start of any new work session where no specific task is
  given. Also triggers on "what should I focus on today" or
  "give me my priorities". Pulls from calendar, priority files,
  and pending items to generate a structured daily brief.
---

# Morning Briefing

## Context Required

Read these files before generating:
- my-priorities.md (your current goals and focus areas)
- current-projects.md (active projects and status)
- Calendar: today's events and attendees
- Pending items from previous sessions

## Steps

1. Read all context files listed above
```

👥 board-of-advisors/SKILL.md PREFERENCE ▲

```
---
name: board-of-advisors
description: >
  Use when the user says "what would my board say", "get
  perspectives on this", "review this from multiple angles",
  "360 review", "what would [role] think about this", or
  asks for feedback from different viewpoints. Also triggers
  on "I need a second opinion" or "who would disagree with
  this". Simulates multi-perspective review from expert
  archetypes, each with defined lenses, biases, and blind
  spots.
---

# Board of Advisors

## Advisor Profiles

Each advisor has: a defined lens, known biases, and
documented blind spots. Read profiles from
advisors/ folder if available, otherwise use defaults:

- The Strategist: long-term positioning, competitive
  dynamics. Bias: overvalues optionality.
- The Operator: execution feasibility, resource
```

📦 All 4 skill folders (Research, Devil's Advocate, Morning Briefing, Board of Advisors) available in the **Expansion Pack → Blueprints** section below.

4 STRATEGIST

Advanced Patterns & System Thinking

ADVANCED PATTERNS

PATTERN	WHAT IT DOES	EXAMPLE

Skill Dispatcher	Meta-skill that routes requests to the right skill	Essential when your library grows past 10-15 skills
Skill Chaining	Output of one skill → input to another	Research → Devil's Advocate → Executive Summary
Loop Skills	Iterative: check → act → check again	Campaign optimization, competitive monitoring, content QA
Agentic Loops	Spawn sub-agents, maintain state, orchestrate	Deep research with parallel workers and aggregation

LOOP SKILLS FOR KNOWLEDGE WORKERS

- **Marketing campaign optimization:** Monitor ad performance → adjust bid strategy → re-check → flag when ROAS drops
- **Competitive intelligence:** Scan competitor sites daily → compare to baseline → alert on meaningful changes
- **Content publishing:** Draft → review against style guide → revise → check compliance → publish
- **Recruitment pipeline:** Monitor candidates → send follow-ups → escalate when stale

TESTING & OPTIMIZATION

> **THE LITMUS TEST**
> If you find yourself having to iterate and refine the output AFTER the skill runs — editing, correcting, restructuring — then the skill itself needs improvement. A well-built skill produces output you can use directly.

- What works on Claude 4.6 might behave differently on Claude 5. Test after model updates.
- A/B test skill versions: is v2 actually better, or just different?
- **Eval proportionality:** Personal writing skill? Quick spot-check. Customer-facing compliance skill? Formal evaluation with edge cases.

WHEN TO RE-EVALUATE YOUR SKILLS

TRIGGER	WHAT TO DO

MODEL CHANGE	New model drops — your gotcha section might be solving problems the new model doesn't have
TOOL CHANGE	Moving between tools? Skills are portable but behaviors differ. Validate.
RESULTS DEGRADE	Before blaming the model: did YOUR context go stale?
BEFORE SCALING	About to share with 50 people? Run proper evals first. Treat it like an AI product.
QUARTERLY	Even if nothing seems broken. Capability skills may have been surpassed by the base model.

5 ARCHITECT

From Personal Skills to Organizational Skill Libraries

THE OPPORTUNITY

Skills are the **pipe dream of every knowledge manager** finally becoming real. Standardization + execution + knowledge — bundled into a single portable artifact that both humans can read and AI can execute.

Organizations doing **skill hackathons** and maintaining shared libraries are seeing massive uplift. Those that aren't? Their people reinvent the wheel in every AI conversation.

THE PLUGIN MODEL

The future = **plugins**: skills + MCP connections bundled. A "sales pipeline review" plugin = the review skill + CRM connection + analytics access + your VP's expected output format.

THE 5-STAGE PROCESS

1 DISCOVERY

Run work audits. Where do people repeat instructions to AI? Survey champions — they've already built skills for themselves. Map the shadow skill landscape.

2

CURATION

Prioritize by frequency × impact × standardization value. Designate skill owners — SMEs, not engineers. The best preference skills are built by the people who do the work.

3

VALIDATION

Test against real scenarios. A/B test vs. unstructured prompting. Get feedback from actual users. Test across tools.

4

BUNDLING

Package into plugins. Define org-wide vs. team vs. personal. Create onboarding bundles. Version control everything.

5

OWNERSHIP & MAINTENANCE

Champions per domain. Quarterly review. Usage tracking. Feedback loops. Deprecation process for obsolete skills.

🎮 EXPANSION PACK

Complete all 5 levels to unlock • or click to peek

SKILL POWER 0%

© 2026 AIDB Training