

The Bank AI Use-Case Inventory Template

Reference card · For community banks and credit unions

A fillable register for tracking AI tools, data, owners, human review, risk tier, vendor controls, evidence, approval status, and re-review cadence.

Use one row per AI use case. A use case is the workflow, not just the tool name. If one vendor supports three workflows, record three rows.

Intake questions

Ask these before opening a row:

- What business problem does this solve?
- What tool or vendor will be used, and is it institution-provisioned?
- What data enters the tool?
- Does customer/NPI, regulated, examination-sensitive, privileged, or security control information enter the tool?
- What output is produced?
- Could the output affect a customer, examiner, control, report, or regulated decision?
- Who signs off before use?
- What evidence will be retained?

Fillable register fields

Build the register as a spreadsheet or table with these columns.

Column	What to fill in
Use Case ID	Stable ID such as AI-OPS-001 or AI-LEND-003.
Use Case Name	Plain-English workflow name, not the product name.

Column	What to fill in
Business Purpose	The operational, risk, compliance, or customer reason this use case exists.
Department	Business line or function that owns the workflow.
Named Owner	One accountable person, not a committee.
Status	Proposed, intake review, sandbox, approved, restricted, blocked, or retired.
Tool/Vendor	Product, vendor, version if known, and whether the tool is public, enterprise, private, or embedded.
Deployment Type	Public AI, approved enterprise account, vendor feature, private deployment, or internal model.
Approved Tool Status	Approved, pending review, exception, prohibited, or not yet assessed.
Data Class	Public, internal, confidential/NPI, regulated, examination-sensitive, or privileged.
Customer/NPI Involved	Yes, no, de-identified, tokenized, or approved private environment only.
Regulated Process Involved	Credit, fraud, BSA/AML, sanctions, complaints, regulatory reporting, HR, legal, or none.
Output Type	Draft, summary, classification, recommendation, score, customer-facing content, report, code, or action.
Human Reviewer	Named role that reviews before use, plus whether review is mandatory or sampled.
Risk Tier	Low, medium, high, or blocked using the guide below.
Required Controls	Approved tool, data limit, human review, validation, monitoring, recordkeeping, escalation, or other controls.
Evidence Retained	Prompt, output, reviewer note, ticket, approval, validation memo, vendor review, or incident link.
Approval Date	Date approved for sandbox or production use.
Next Review Date	Quarterly for high-risk rows; at least annually for medium/low rows.
Re-review Trigger	Tool change, data class change, policy change, vendor update, incident, complaint, or move to production.

Column	What to fill in
Incident/Exception Link	Ticket, exception log, risk acceptance, or incident record if applicable.
Vendor Review Link	Third-party due diligence, contract review, SOC report, or vendor monitoring file.

Vendor lifecycle fields

Use these fields whenever the row involves a vendor, embedded vendor feature, hosted model, or third-party AI capability.

Lifecycle stage	Inventory question
Planning rationale	What business problem does the AI tool solve, and why is a third party needed?
Due diligence status	Has InfoSec, Compliance, Risk, and the business owner reviewed the tool for this use?
Contract review status	Does the agreement cover confidentiality, audit, regulatory access, breach notice, and use limits?
Data-use terms	May the vendor use prompts, outputs, or bank data for model training or product improvement?
Model-training restriction	Is training on bank/customer data prohibited or opt-out confirmed in writing?
Retention/deletion terms	How long are prompts and outputs retained, and how are they deleted?
Subcontractor/fourth-party review	What subprocessors, hosted models, or infrastructure providers are involved?
Breach-notice terms	What notice timeline and evidence preservation requirements apply?
Monitoring owner	Who reviews performance, incidents, complaints, and vendor changes?
Next vendor review date	When is the next monitoring review due?
Termination/data-return plan	How will access be revoked and bank data returned or deleted?

Risk-tier guide

Tier on the highest factor that applies. The purpose is consistency across departments, not false precision.

Tier	Definition
Low	Internal drafting, public data, no customer impact, no regulated decision, approved tool, and human review before use.
Medium	Internal process support, customer-facing draft content, confidential internal data, or operational workflow support where human review is required.
High	Decision support for credit, fraud, BSA/AML, sanctions, complaints, regulatory reporting, customer-impacting workflows, or NPI used only in approved private environments.
Blocked	Public AI tool with NPI, SAR/AML detail, examination-sensitive information, privileged material, security controls, or final regulated decisions.

Model-risk note: use current model-risk guidance and principles, including SR 26-2 where applicable, for AI use cases that inform quantitative, customer-impacting, or regulated decisions. Do not imply every generative AI use case at every institution is automatically subject to the same model-risk regime.

Approval workflow

Move each row through this path:

1. Proposed - department submits intake questions and draft row.
2. Intake review - AI Program Owner confirms the row is a use case and names an owner.
3. Data classification - InfoSec or data owner confirms the highest data class and prohibited inputs.
4. Vendor/tool review - TPRM, InfoSec, and Compliance confirm approved-tool and contract status.
5. Risk tier assigned - Compliance/Risk assigns low, medium, high, or blocked.
6. Human review control defined - owner names reviewer, review standard, and evidence retained.
7. Approved, restricted, or blocked - approval is recorded with date, conditions, and next review date.
8. Quarterly or annual re-review - cadence follows the tier and any trigger event.

Sample rows

Use these to calibrate the first pass.

Example	Use case	Data/tool	Tier	Required control
Green/Low	Summarize public regulator press releases for internal training.	Public data in an approved tool.	Low	Human editor confirms accuracy before training use.
Yellow/Medium	Draft customer email language using approved templates and no customer data.	Internal templates in an approved enterprise tool.	Medium	Marketing and Compliance review before sending.
Red/Blocked	Enter loan-file details or SAR investigation notes into a public AI tool.	NPI, loan file, SAR/AML, or examination-sensitive content in a public tool.	Blocked	Do not use. Escalate to AI Program Owner and Compliance if attempted.

Maintenance cadence

- High-risk rows: review at least quarterly and whenever the tool, model, vendor, data class, workflow, or applicable guidance changes.
- Medium and low rows: review at least annually and on any material trigger event.
- Board or committee reporting: summarize count of use cases by tier, new rows, blocked rows, open exceptions, incidents, and overdue reviews.
- Retired rows: keep enough history to show when use stopped, who approved retirement, and where records were retained.

Next step: AI Use-Case Inventory Workshop

Use this template to run a 60-minute workshop that identifies shadow AI use, classifies risk tiers, assigns owners, and produces the first board-ready inventory.

Bring Compliance, Risk, InfoSec, Operations, Lending, BSA/AML, Marketing, and the AI Program Owner. Start with tools staff already use, then ask where vendors have added AI features inside existing systems. Leave with at least three completed rows, blocked uses named, owners assigned, and review dates scheduled.

Adapt before adoption

This template is a starting point, not legal or compliance advice. Adapt tier definitions, approval roles, cadence, data classes, and vendor-control fields to your institution's size, complexity, charter, regulator, data standards, and third-party risk program before adoption.

Sources

- AIEOG AI Lexicon and Financial Services AI Risk Management Framework, US Treasury / FBIIC / FSSCC, February 2026
- SR 26-2, Revised Guidance on Model Risk Management, Federal Reserve / OCC / FDIC, April 17, 2026
- Interagency Guidance on Third-Party Relationships: Risk Management, Federal Reserve / OCC / FDIC, June 2023
- NIST AI Risk Management Framework (AI RMF 1.0)
- ECOA / Regulation B and CFPB adverse-action guidance for credit-related use cases