[Ai] Banking Institute

AI Use-Case Inventory

Reference card · For community banks and credit unions

The single register of every place AI touches your institution: what it does, who owns it, what data it sees, and how closely a human watches it. Maintain it. Bring it to your AI committee. Hand it to your examiner before they ask — because they will ask for this one first.

Why this is the artifact examiners ask for first

There is still no comprehensive AI-specific banking law. Federal financial regulators supervise AI through existing frameworks — model risk management, third-party risk management, fair lending, BSA/AML, and consumer protection — applied through risk-based examinations (GAO-25-107197, May 2025). That has a direct consequence: an examiner cannot ask "are you following the AI rule," so they ask the next best question — "show me everything you're doing with AI and how you govern it." The inventory is that answer in one document.

It is also a baseline expectation in its own right. The AIEOG AI Lexicon (US Treasury / FBIIC / FSSCC) names the "AI use case inventory" as a foundational governance practice for institutions of any size. And model risk management guidance has long expected a complete inventory of models in use, under development, or recently retired. The revised interagency guidance, SR 26-2 (Federal Reserve / OCC / FDIC, April 17, 2026), superseded SR 11-7 and SR 21-8 and made the framework explicitly risk-based and tailored to an institution's size and model risk profile. It is described as most relevant to banks over $30 billion in assets, but the agencies note it can still apply to smaller institutions where model use is significant — and supervisory focus expressly includes whether the inventory captures all material models and third-party tools. For a community bank or credit union, the inventory is the cheapest, highest-leverage examiner-readiness move you can make.

If you do nothing else with AI governance this year, keep this list current.

The column schema

One row per use case. These are the fields an examiner expects to see, each with a one-line definition.

Column	What goes in it
Use case	The actual task in plain language — what the AI does, not the product name (e.g., "draft BSA SAR narratives").
Owner	A named person accountable for this use case. Not "the committee," not a department.
Business purpose	Why the institution does this — the outcome it serves (efficiency, service, detection, compliance).
Tool & vendor	The specific product, version, and vendor (e.g., "Microsoft 365 Copilot, Microsoft"). Note "in-house" if built internally.
Data classes touched	The data the use case can see: Green (public), Yellow (internal, non-NPI), Red (NPI / PII / loan files).
Risk tier	Low / Medium / High, assigned with the rubric below.
Human-in-the-loop?	Yes / No — whether a person reviews or approves output before it has effect. For "Yes," note the control (mandatory / sampled).

Column	What goes in it
Regulatory touchpoints	The rules the use case implicates (e.g., ECOA/Reg B, BSA/AML, GLBA, UDAAP, fair lending, Reg E).
Validation status	Validated / In review / Not validated / Vendor-attested — the state of independent testing of the tool for this use.
Review cadence	How often this row is re-examined (e.g., quarterly, annually). Higher tier means more frequent.
Last reviewed	The date a human last verified the row still describes how the use case actually runs (YYYY-MM-DD).
Status	Production / Pilot / Proposed / Retired.

Risk-tiering rubric

Tier on the highest factor that applies — if any single row pushes High, the use case is High. A banker should be able to tier a use case in 30 seconds.

Factor	Low	Medium	High
Data touched	Green only	Yellow	Red (NPI / PII / loan or member files)
Effect of output	Internal draft, human rewrites fully	Customer-facing after review; operational input	Influences a credit, account, BSA/AML, or adverse-action decision
Autonomy	Output always rewritten by a person	Output edited/approved by a person	Output can act or post with no human check
Regulated domain	None	Indirect (marketing, internal ops)	Lending, BSA/AML, deposits, collections, fair lending
Customer impact if wrong	Negligible	Recoverable inconvenience	Financial, legal, or fair-lending harm

Rule of thumb: anything that touches Red data, or sits in or near a credit / account / BSA-AML decision, is **High** and should run with a human-in-the-loop. Marketing copy on public facts with a human editor is typically **Low**.

Filled example inventory

A worked example across departments. Use it to see exactly how to complete a row. Dates and vendors are illustrative.

Use case	Owner	Business purpose	Tool & vendor	Data classes	Risk tier	HITL?	Regulatory touchpoints	Validation status	C
Pre-screen consumer loan applications for completeness and route to underwriters	Dana Whitfield, VP Lending	Faster, more consistent intake triage	Underwriting Assist, LendingCoVendor	Red	High	Yes (mandatory — underwriter reviews every routing)	ECOA / Reg B, fair lending, FCRA	In review	C
Draft BSA/AML SAR narratives from analyst-supplied facts	Marcus Reyes, BSA Officer	Reduce narrative drafting time, improve consistency	In-house GPT on approved tenant	Red	High	Yes (mandatory — BSA Officer edits and signs)	BSA / AML, SAR confidentiality	Validated	C

Use case	Owner	Business purpose	Tool & vendor	Data classes	Risk tier	HITL?	Regulatory touchpoints	Validation status
Summarize internal policy documents for staff training	Priya Nair, L&D Manager	Speed up training content development	Microsoft 365 Copilot, Microsoft	Yellow	Medium	Yes (sampled — reviewer spot-checks)	Internal policy accuracy	Vendor-attested
Generate marketing email copy on public product features	Sam Ortiz, Marketing Director	Faster campaign drafting	Jasper, Jasper AI	Green	Low	Yes (mandatory — editor reviews before send)	UDAAP, advertising / Reg DD claims	Not validated
Classify and route inbound member service emails	Lee Tran, Ops Manager	Reduce response time, route to right queue	ServiceRouter AI, OpsVendor	Yellow	Medium	No (auto-routes; misroutes corrected downstream)	GLBA (data handling), Reg E (if disputes)	In review
Chatbot answering general account FAQs on the website	Jordan Kim, Digital Banking	Self-service for common questions	ChatAssist, FinChatVendor	Yellow	High	Yes (escalates to human; no transactional actions)	UDAAP, GLBA, Reg E	In review

How to stand it up and maintain it

1. **Assign one owner of the inventory itself** — typically a named officer in risk or compliance. The inventory needs a person, not a committee, on the hook for keeping it current.
2. **Survey every department.** Ask the same two questions everywhere: "Are you using AI for any work task?" and "Is any vendor giving you AI-powered features you may not have classified as AI?" The second question is how you catch shadow AI buried in tools you already license.
3. **Capture one row per use case**, even when one tool covers several use cases. Tier each row with the rubric. Most community institutions surface 9–14 use cases on the first pass.
4. **Set cadence by tier.** Review High-tier rows quarterly, Medium and Low at least annually. Re-review any row immediately when the vendor materially updates the tool, the data classes change, or the use case moves from pilot to production.
5. **Feed board reporting.** Each cycle, roll the inventory into a short summary for the board or its risk committee: count of use cases by tier, what's new since last period, what's in pilot, any High-tier row without completed validation, and any open issues. The inventory is the source of record; the board summary is its quarterly read-out.

Common mistakes

- **"The committee" as owner.** Accountability has to land on a person. A committee can govern; it cannot be the owner of a use case.
- **Shadow AI not captured.** The biggest gaps are AI features inside tools you already use — CRM, core, productivity suites, fraud tools. If you only list tools you bought "for AI," the inventory is incomplete.
- **No last-reviewed date, or a stale one.** An undated row is an unverified row. Examiners read the dates first.
- **Listing tools instead of use cases.** "Microsoft 365 Copilot" is not a row. "Draft BSA SAR narratives in Copilot" is. One tool can generate many rows, each with its own tier and oversight.
- **Treating pilots as off-inventory.** Pilots touching Red data are exactly what examiners want to see governed. Log them with status = Pilot.
- **No human-in-the-loop on High-tier rows.** Any use case influencing a credit, account, or BSA/AML decision should show how a person checks it.

- **One-and-done.** An inventory built for an exam and never updated is worse than none — it documents that you know what you have and stopped watching.

Adapt before adoption

This template is a starting point, not a policy. Tier definitions, cadence, data classes, and ownership should be adapted to your institution's size, risk profile, charter, and regulator before adoption. This artifact is general information, not legal or compliance advice; confirm specifics with your own counsel and compliance function.

Sources

- AIEOG AI Lexicon, US Treasury / FBIIC / FSSCC
- SR 26-2, Revised Guidance on Model Risk Management, Federal Reserve / OCC / FDIC, April 17, 2026 (supersedes SR 11-7, April 4, 2011, and SR 21-8)
- Interagency Guidance on Third-Party Relationships: Risk Management, Federal Reserve / OCC / FDIC, June 2023
- GAO-25-107197, Artificial Intelligence: Use and Oversight in Financial Services, US GAO, May 2025
- Equal Credit Opportunity Act (15 U.S.C. § 1691) and Regulation B (12 CFR Part 1002), CFPB
- Gramm-Leach-Bliley Act Privacy Rule (Title V, Subtitle A)