

Fair-Lending AI Review Checklist

Reference card · For community banks and credit unions

For institutions that have, or are about to have, AI in or near a credit decision. Run this before deployment, then again on a recurring cadence.

Why this exists

ECOA and Regulation B do not have an "AI exception." Disparate-impact obligations attach to any process that influences a credit decision, whether the influence comes from a human, a static model, or an AI tool. The institutions that run this review before being asked are positioned to defend the program; the ones that don't are positioned to explain why.

Scope this review

Identify every AI-assisted process that:

- Generates, recommends, or filters credit decisions
- Sets pricing tiers, terms, or limits
- Determines marketing eligibility for credit products
- Triagees applicants for human review
- Drafts adverse-action reasoning or supporting documentation

Each one needs its own review. The same process may need separate reviews for each product line.

Pre-deployment checklist

DATA AND DESIGN

- ☐ Training or input data documented, including the date range, source systems, and any known gaps

- ☐ Protected-class indicators identified in the data layer (race, sex, age, marital status, national origin, disability, familial status, receipt of public assistance, exercise of ECOA rights)
- ☐ Proxy variables flagged (ZIP code, surname, neighborhood markers that may correlate with protected class)
- ☐ Decisions documented end-to-end: what data goes in, what comes out, who acts on it

DISPARATE-IMPACT TESTING

- ☐ Approval-rate gap measured across protected classes
- ☐ Score / recommendation distribution compared across protected classes
- ☐ Outcome variance (loan performance, default, prepay) compared across protected classes
- ☐ Results compared to the non-AI baseline (same process without the AI tool)
- ☐ Documented threshold for what constitutes a material gap, and the action triggered at the threshold

EXPLAINABILITY

- ☐ Adverse-action reasoning can be produced for any AI-assisted decision in a form that satisfies Regulation B
- ☐ Principal reasons map to specific input factors, not opaque scores
- ☐ Documentation is reviewable by a compliance officer without vendor assistance

GOVERNANCE

- ☐ Named owner for the AI-assisted process
- ☐ Named compliance reviewer
- ☐ Documented oversight level (automated / sampled / mandatory)
- ☐ Escalation path when a metric crosses an internal threshold
- ☐ Logged in the institution's AI use-case inventory

Recurring review checklist

Run quarterly at minimum. Annually if the process is low-volume.

- ☐ Re-pull disparate-impact metrics
- ☐ Compare to the prior period — flag any worsening trend
- ☐ Verify the AI tool has not been materially updated by the vendor since the last review; if it has, re-run the full pre-deployment review

- ☐ Review adverse-action samples for explainability
- ☐ Report results to the same forum that receives the institution's standard fair-lending program reporting

Decisions log

Maintain a short log of decisions made under this review: dates, who signed off, what changed. The log is what an examiner reads first.

Citations

- Equal Credit Opportunity Act (15 U.S.C. § 1691)
- Regulation B (12 CFR Part 1002), CFPB
- SR 11-7 Guidance on Model Risk Management, Federal Reserve / OCC
- AIEOG AI Lexicon, US Treasury / FBIIC / FSSCC, February 2026
- GAO-25-107197, US GAO, May 2025