The AI Banking Institute
Turning Bankers into Builders

ROLE PLAYBOOK · BSA / AML · V2.0

The BSA / AML Officer's AI-Native Playbook

Sharper narratives. Cleaner documentation. Safer automation. A practical operating playbook for BSA/AML leaders at community banks and credit unions who want to use AI without weakening judgment, violating data boundaries, or creating unexplainable compliance risk.

The operating rule

AI prepares. Humans decide. Evidence proves.

AI can help analysts draft, structure, summarize, compare, train, and document. It cannot own the suspicious-activity judgment, final SAR decision, sanctions disposition, investigation closure, or regulatory accountability.

WHAT'S INSIDE

Who this is for

- BSA / AML Officers
- Compliance Officers overseeing financial-crime controls
- AML Analysts and team leads
- Operations leaders supporting alert disposition and case documentation
- InfoSec / vendor-risk partners evaluating approved AI tools
- Executives asking where AI can safely help the BSA function

The problem

BSA/AML work is overloaded with manual drafting, inconsistent documentation, alert noise, analyst training burden, and institutional knowledge trapped in senior staff members' heads. AI can reduce that burden, but BSA/AML has one of the hardest data boundaries in the bank. SAR information, customer NPI, transaction-level detail, case notes, and investigation records cannot be casually pasted into public AI tools.

The opportunity

Used correctly, AI can help BSA/AML teams improve SAR narrative structure, standardize CDD baseline language, generate synthetic training scenarios, draft procedure documentation, summarize typologies and false-positive patterns, prepare alert-tuning conversations, build repeatable review checklists, and document evidence consistently.

What this playbook gives you

1. A clear data-boundary model for AI use.
2. Five practical BSA/AML plays.
3. Copy/paste prompts and templates.
4. Human review checkpoints.
5. A 30-day implementation plan.
6. Evidence packet requirements.
7. Metrics to prove value without weakening control.

BSA/AML lives behind a hard line

BSA/AML is not a casual AI experimentation area. It includes customer data, restricted SAR information, law-enforcement-sensitive workflows, confidential investigation notes, and high-stakes regulatory obligations.

> **THE RED LINE**
>
> No customer name, account number, transaction detail, SAR narrative, case note, alert detail, or investigation record goes into a public AI tool. The plays in this document are designed for de-identified patterns, synthetic scenarios, templates, procedure drafts, aggregated metrics, or approved enterprise environments with documented data-handling terms.

What AI is genuinely good for in BSA/AML

- **Pattern description** — structuring, layering, funnel accounts, unusual cash activity, business-purpose mismatch, rapid movement of funds.
- **Narrative scaffolding** — organizing analyst notes into the who / what / when / where / why / how structure.
- **CDD baseline documentation** — describing expected activity for customer segments in plain language.
- **Procedure writing** — turning tacit senior-officer knowledge into written SOPs.
- **Training scenarios** — synthetic case patterns for analyst practice.
- **Aggregate alert analysis** — summarizing false-positive patterns at the typology level.

AI should not independently decide

- Whether to file a SAR.
- Whether activity is suspicious.
- Whether an alert should be closed.
- Whether a sanctions match is a true match.
- Whether an investigation is complete.
- Whether a customer's risk rating changes.

Five layers of a control system

A modern BSA/AML playbook is not just a procedural binder for CDD, EDD, monitoring, SARs, and CTRs. It is a control system that defines what AI may prepare, what humans must decide, what evidence must be retained, what data may be used, which tools are approved, which outputs require review, and how the institution proves the workflow is working.

Layer 1 — Data Boundary

Public · Internal · Confidential · Customer NPI · SAR-restricted · Law-enforcement-sensitive. Public AI tools stay in the Public / generic / synthetic / redacted lane only.

Layer 2 — AI Workbench

AI may assist with drafting, summarizing, pattern description, training scenario generation, typology research, and procedure structuring.

Layer 3 — Human Decision Gates

A qualified human must decide or sign off on: SAR filing, final SAR narratives, alert closure, sanctions disposition, investigation closure, risk-rating changes, and escalation to law enforcement or senior management.

Layer 4 — Validation and Monitoring

Track output quality, reviewer changes, false-positive themes, analyst overrides, escalations, exceptions, rework, and drift in alert patterns or typology language.

Layer 5 — Evidence and Explainability

Retain source material, prompt, AI output, reviewer name, review changes, final decision, timestamp, and retention location. If the institution cannot reconstruct why an alert, narrative, or output was accepted, changed, escalated, or rejected, the workflow is not ready.

AI can prepare. Humans must decide.

WORKFLOW	AI CAN PREPARE	HUMAN MUST DECIDE	EVIDENCE RETAINED
SAR narrative	Structure, draft sections, missing-fact flags	Whether to file and final wording	Draft, reviewer edits, source notes, final support
Sanctions alert	Entity research summary, false-positive rationale draft	Ambiguous match disposition	Sources reviewed, reviewer, disposition
KYB / UBO	Ownership structure summary from approved sources	Beneficial ownership conclusion	Source docs, reviewer, final profile
CDD baseline	Segment baseline description	Customer-specific risk rating	Baseline, reviewer, approval date
Alert tuning	Pattern summary, false-positive themes	Threshold or scenario changes	Aggregate inputs, tuning recommendation, final decision
Analyst training	Synthetic case narratives	Training evaluation and analyst readiness	Scenario, answer key, trainer notes
Procedure writing	SOP draft from approved process notes	Final procedure approval	Draft, approver, published SOP

Three lanes for BSA/AML AI use

Green — generally safe with synthetic or public information

- Generate synthetic training scenarios.
- Draft internal typology explainers.
- Summarize public FinCEN guidance.
- Create onboarding quizzes.
- Build blank templates.
- Draft SOP structure without customer facts.

Rules: no customer data; no SAR content; no investigation details; human review before distribution.

Yellow — approved tool only, with redaction and review

- Draft CDD segment baseline descriptions.
- Structure redacted analyst notes.
- Summarize aggregate alert patterns.
- Draft internal procedures.
- Create false-positive pattern summaries from aggregate data.

Rules: approved tool only; strip customer identifiers; no transaction-level details in public tools; human review required; retain reviewed version only.

Red — restricted or prohibited except in approved enterprise workflow

- SAR narratives using real case facts.
- Alert disposition using customer-level data.
- Sanctions match decisions.
- Investigation closure.
- Customer risk-rating updates.
- Transaction-monitoring model changes.

Rules: no public AI tools; approved enterprise environment only; documented data-handling terms; human sign-off; evidence retained; potential model-risk / vendor-risk review depending on use.

Pick one safe first play

Start with one of these:

1. CDD segment baseline builder.
2. Synthetic analyst training scenario.
3. Procedure-to-SOP draft.
4. SAR narrative structure checklist using fictional facts.
5. Aggregate false-positive pattern summary.

The five-step pattern

1. Pick the workflow.
2. Confirm the data boundary.
3. Use the provided prompt or template.
4. Review the output with a qualified human.
5. Save the reviewed artifact and note the retention location.

DO NOT START WITH

- Real SAR narratives.
- Real alert detail.
- Customer transaction data.
- Sanctions match decisions.
- Case closure recommendations.
- Customer risk scoring.

PLAY 1

Better SAR narrative scaffolding

Help analysts organize SAR narratives around who, what, when, where, why, and how — without owning the suspicious-activity judgment.

Why it matters

SAR narratives are often inconsistent because they are written by analysts with different experience levels under time pressure. AI can help standardize structure without owning the suspicious-activity judgment.

Approved inputs

- **Public tools:** synthetic case facts, training examples, blank SAR narrative template, fictional typology notes.
- **Approved enterprise tools only:** redacted analyst notes, approved internal templates, controlled case-support notes if permitted by tool terms and policy.

Prohibited inputs

- Customer names.
- Account numbers.
- Transaction-level detail.
- SAR narratives from real filings.
- Case management exports.
- Confidential investigation notes in public tools.

PROMPT TEMPLATE · SAR NARRATIVE SCAFFOLD

```
You are a BSA/AML quality reviewer helping structure a draft SAR narrative.

Use only the facts provided below. Do not invent facts, dates, names, account activity,
transaction amounts, or suspicious-activity conclusions.

Organize the narrative into:
  1. Who is involved
  2. What activity occurred
  3. When the activity occurred
  4. Where the activity occurred
  5. Why the activity is unusual
  6. How the activity fits the typology
  7. Missing facts to verify
  8. Reviewer checklist

Label the output as draft support only. The analyst owns the final SAR narrative and filing
decision.

Facts:
[INSERT FICTIONAL OR APPROVED / REDACTED FACTS]
```

Expected output

- Draft narrative structure.
- Missing-fact list.
- Analyst review checklist.
- Typology alignment notes.

Human review & evidence

- BSA analyst for first review.
- BSA officer or designated reviewer for final narrative approval.
- Retain: input facts, scaffold output, analyst edits, reviewer sign-off, final narrative support location.

Metrics to track

- Draft time per narrative.
- Reviewer rework rate.
- Missing-fact flags per draft.
- Narrative quality review results.

PLAY 2

De-noise the alert queue

Summarize false-positive patterns at the typology level to prepare for alert-tuning conversations.

Why it matters

Analysts often know which alerts are noisy, but that knowledge sits in frustration, not structured evidence. AI can help convert aggregate patterns into a cleaner tuning conversation.

Approved inputs

- Aggregate alert counts.
- Typology labels.
- False-positive reason categories.
- Quarter-over-quarter trend summaries.
- Analyst feedback themes without customer detail.

Prohibited inputs

- Individual alert records.
- Customer names.
- Transaction details.
- Case notes.
- SAR references.
- Investigation files.

PROMPT TEMPLATE · FALSE-POSITIVE PATTERN SUMMARY

```
You are assisting a BSA officer preparing for a quarterly alert-tuning discussion.

Use only the aggregate pattern data below. Do not make a tuning decision. Summarize the major
false-positive themes and prepare questions for the AML vendor or model owner.

Output:
  1. Top alert types by false-positive volume
  2. Pattern description for each alert type
  3. Likely source of noise
  4. Questions for vendor / model owner
  5. Data needed before any tuning decision
  6. Human decision checklist

Aggregate data:
[INSERT AGGREGATE COUNTS AND PATTERN DESCRIPTIONS]
```

Expected output

- Quarterly false-positive summary.
- Tuning discussion agenda.
- Questions for vendor / model owner.
- Human decision checklist.

Human review & evidence

- BSA officer reviews and determines whether any tuning action is warranted.
- Retain: aggregate input data, AI summary, vendor discussion notes, final tuning or no-change decision.

Metrics to track

- False-positive rate by alert type.
- Analyst hours spent on top noisy alerts.
- Number of vendor questions resolved.
- Tuning decision cycle time.

PLAY 3

CDD baseline builder

Draft plain-English expected-activity descriptions for customer segments.

Why it matters

Senior analysts often understand what normal looks like for customer segments, but that knowledge is not always documented in files or onboarding material.

Approved inputs

- Generic segment descriptions.
- Product types.
- Business type.
- Expected activity pattern by segment.
- Public business description.
- Senior analyst notes without customer identifiers.

Prohibited inputs

- Customer name.
- Account number.
- Actual transaction-level activity.
- Real balances.
- SAR or alert notes.

PROMPT TEMPLATE · CDD SEGMENT BASELINE

```
You are helping a BSA officer draft a plain-English CDD baseline for a customer segment.

Do not describe a specific customer. Do not include customer names, account numbers,
balances, or transaction-level activity.

Segment:
[SEGMENT NAME]

Known expected activity:
[GENERIC ACTIVITY PATTERN]

Output:
  1. One-paragraph expected activity baseline
  2. Common variations that may still be normal
  3. Activity that may require review
  4. Questions an analyst should ask when activity differs
  5. Reviewer checklist

Label this as a draft baseline requiring BSA officer approval.
```

Expected output

- One-paragraph CDD baseline.
- Common variations.
- Escalation triggers.
- Analyst review questions.

Human review & evidence

- Senior BSA officer approves each baseline before use.
- Retain: segment definition, draft baseline, reviewer approval, date added to library.

Metrics to track

- Number of segment baselines created.
- Analyst onboarding time.
- File documentation consistency.
- Reviewer corrections per baseline.

PLAY 4

Synthetic training scenario generator

Create synthetic BSA/AML practice cases for analyst training.

Why it matters

Analysts need reps. Real cases are sensitive. Synthetic cases let analysts practice typology recognition without exposing customer data.

Approved inputs

- Typology name.
- Difficulty level.
- Training objective.
- Fictional customer profile.
- Fictional activity pattern.

Prohibited inputs

- Real cases.
- Real SARs.
- Real alert details.
- Customer identifiers.
- Case management records.

PROMPT TEMPLATE · SYNTHETIC BSA TRAINING SCENARIO

```
Create a synthetic BSA/AML training scenario for analyst practice.

Typology:
[STRUCTURING / LAYERING / FUNNEL ACCOUNT / BUSINESS-PURPOSE MISMATCH / OTHER]

Difficulty:
[BEGINNER / INTERMEDIATE / ADVANCED]

Requirements:
  - Use fictional names and fictional activity only
  - Include enough details for an analyst to identify the typology
  - Do not include real customer data
  - Provide an answer key for the trainer
  - Include three discussion questions
  - Include one misleading but plausible fact that should not drive the conclusion

Output:
  1. Scenario narrative
  2. Analyst task
  3. Expected observations
  4. Trainer answer key
  5. Discussion questions
  6. Risk notes
```

Expected output

- Synthetic case narrative.
- Analyst exercise.
- Answer key.
- Discussion questions.

Human review & evidence

- Trainer or BSA officer reviews before use.
- Retain: scenario prompt, scenario output, trainer answer key, training date, analyst completion notes.

Metrics to track

- Number of training scenarios created.
- Analyst completion rate.
- Assessment scores.
- Typology recognition improvement.

PLAY 5

Procedure writing on demand

Turn senior officer knowledge into written SOPs.

Why it matters

BSA/AML procedures often live in the head of the senior officer. AI can structure dictated process knowledge into a reviewable SOP.

Approved inputs

- Process transcript.
- Approved procedure notes.
- Internal process outline without customer details.
- Review owner.
- Escalation rules.

Prohibited inputs

- Customer records.
- Real case facts.
- SAR information.
- Investigation notes.
- Confidential examiner correspondence in public tools.

PROMPT TEMPLATE · BSA PROCEDURE-TO-SOP

```
You are helping a BSA officer convert process notes into a written SOP.

Do not add policy requirements that are not in the notes. Flag missing information with
[VERIFY]. Do not include customer data or case details.

Process notes:
[INSERT APPROVED PROCESS NOTES]

Output:
  1. Purpose
  2. When to use this procedure
  3. Roles and responsibilities
  4. Step-by-step process
  5. Escalation triggers
  6. Required documentation
  7. Review checklist
  8. Retention guidance
  9. Open questions marked [VERIFY]

Label the output as a draft SOP requiring BSA officer and compliance review.
```

Expected output

- Draft SOP.
- Open-question list.
- Review checklist.
- Retention guidance.

Human review & evidence

- BSA officer and compliance officer approve before publication.
- Retain: source notes, draft SOP, reviewer edits, final approved SOP, publication date.

Metrics to track

- SOPs drafted.
- SOPs approved.
- Review cycle time.
- Training use.
- Exceptions reduced.

Five reusable templates

Template 1 · BSA/AML AI use-case intake form

FIELD	RESPONSE
Use case name	
Business purpose	
BSA/AML process affected	
Proposed tool	
Tool approval status	
User role	
Input data class	Public / Internal / Confidential / Customer NPI / SAR-restricted
Output type	Draft / Summary / Checklist / Recommendation / Decision support
Risk tier	Green / Yellow / Red
Human reviewer	
Approval checkpoint	
Retention rule	
Escalation trigger	
Initial reviewer notes	

Template 2 · SAR narrative review checklist

- ☐ The narrative explains who, what, when, where, why, and how.
- ☐ The narrative uses only verified facts.
- ☐ No AI-introduced fact remains unsupported.
- ☐ Missing facts were reviewed or removed.
- ☐ The analyst owns the suspicious-activity judgment.

- ☐ The BSA officer or designated reviewer approved the final language.
- ☐ The support file points to source evidence.
- ☐ The AI output is retained only according to approved procedure.

Template 3 · CDD baseline library entry

FIELD	RESPONSE
Segment name	
Segment description	
Expected products	
Expected activity	
Normal variations	
Review triggers	
Analyst questions	
Approved by	
Approval date	
Review cadence	

Template 4 · Synthetic scenario trainer sheet

FIELD	RESPONSE
Scenario name	
Typology	
Difficulty	
Training objective	
Scenario summary	
Analyst task	
Expected observations	
Answer key	
Discussion questions	
Trainer notes	

Template 5 · AI-supported BSA evidence packet

For each AI-supported BSA workflow, retain:

- ☐ Use-case summary.
- ☐ Tool used.
- ☐ Tool approval status.
- ☐ Input data class.
- ☐ Prompt or instruction.
- ☐ Source material used.
- ☐ AI output.
- ☐ Reviewer name.
- ☐ Reviewer changes.
- ☐ Final disposition.
- ☐ Approval date.
- ☐ Retention location.

SECTION 9 · THE 30-DAY IMPLEMENTATION PATH

Week-by-week outcomes

Week 1 — confirm boundaries and pick one play

Outcome: approved enterprise tool status confirmed with Compliance and InfoSec; public AI red-line rules distributed to BSA team; one safe first play selected; data-boundary matrix reviewed.
Recommended first play: CDD baseline builder or synthetic training scenarios.
Evidence: signed red-line rules, use-case intake form, approved input list.

Week 2 — build and test the first artifact

Outcome: first prompt or template drafted; sandbox test completed using synthetic or approved redacted data; human review checklist applied; initial output saved as draft.
Evidence: prompt, draft output, reviewer notes, improvement log.

Week 3 — expand to a second play

Outcome: add a second use case (procedure writing or CDD baseline library); create two to three reviewed artifacts; begin measuring time saved or quality improvement.
Evidence: two reviewed artifacts, measurement tracker, review checklist.

Week 4 — review, decide, and formalize

Outcome: BSA officer reviews pilot outputs; Compliance confirms whether the workflow can continue; InfoSec confirms approved tool / data-boundary status; next month's use cases prioritized.
Evidence: 30-day pilot summary, final approved artifacts, go / no-go / revise decision, next-step roadmap.

SECTION 10 · MEASUREMENT MODEL

Measure one operational improvement per play

PLAY	PRIMARY METRICS
SAR narrative scaffolding	Draft time per narrative · Reviewer correction rate · Missing facts flagged · Narrative quality score
Alert queue patterning	Analyst hours on top false-positive alert types · False-positive patterns identified · Vendor questions resolved · Tuning decision cycle time
CDD baseline builder	Baselines created · Analyst onboarding time · File consistency score · Reviewer corrections
Synthetic training	Scenarios created · Analyst training reps completed · Assessment score improvement · Trainer prep time saved
Procedure writing	SOPs drafted · SOPs approved · Review cycle time · Exceptions reduced

SECTION 11 · EVIDENCE PACKET

What to save for every AI-supported play

Minimum evidence packet

1. Use-case intake form.
2. Data-boundary classification.
3. Prompt or instruction.
4. Input source description.
5. AI output.
6. Human reviewer.
7. Reviewer changes.
8. Final approved version.
9. Retention location.
10. Next review date.

Approved evidence labels

- Draft AI output — not approved.
- Reviewed AI-assisted output.

- Approved SOP.
- Synthetic training scenario.
- Aggregate analysis only.
- Enterprise tool required.
- Public AI prohibited.

One-slide review format

ITEM	STATUS
Use case selected	
Tool approved	
Data class	
Risk tier	
Human reviewer	
Artifacts created	
Time saved	
Quality improvement	
Issues found	
Decision	Continue / revise / stop

Board-level message

> The BSA/AML team piloted AI-supported drafting and documentation using approved data boundaries and human review. AI did not make suspicious-activity decisions, close investigations, or file SARs. The pilot produced reviewed artifacts, measured time savings, and established an evidence-retention pattern for future AI-supported work.

SECTION 13 · WHAT GOOD LOOKS LIKE

The seven questions every BSA team should answer at day 30

1. What AI use cases are allowed?
2. What data is prohibited?
3. What tool may be used?
4. Who reviews the output?
5. What evidence is retained?
6. What metric proves value?
7. Which use cases remain off-limits?

Maturity target

The goal is not full automation. The goal is controlled capability. At the end of the first month, the institution should have one approved AI-assisted BSA workflow, one reviewed artifact, one prompt or template saved for reuse, one human review checklist, one evidence packet, and one measurable result.

SECTION 14 · NEXT STEP

The AiBI path

1. Take the AI Readiness Assessment.
2. Assign the BSA/AML team to the Foundation Course.
3. Complete the Red / Yellow / Green Use Card and Safe AI Use Checklist.
4. Run the BSA/AML synthetic scenario sandbox.
5. Build the first BSA/AML workflow SOP.
6. Save the reviewed asset to the Toolbox.
7. Reassess after 90 days.

> **AI is not the analyst. AI is the workbench. The analyst still investigates. The BSA officer still decides. The institution still owns the filing, the control, and the evidence. The value of AI in BSA/AML is not replacing judgment — it is making judgment easier to document, easier to review, and easier to defend.**

Sources: 31 CFR 1020.320 (SAR information confidentiality); FinCEN SAR narrative guidance; GLBA / Regulation P; Interagency Third-Party Risk Management Guidance; SR 11-7 (Federal Reserve); AIEOG AI Lexicon (US Treasury / FBIIC / FSSCC, February 2026); AiBI Foundations curriculum.