

IN-DEPTH ASSESSMENT FULFILLMENT · V2.1

Your First AI Win — a 90-day playbook

The companion playbook to your In-Depth Readiness Assessment. A 90-day plan that takes the highest-leverage finding from your report and turns it into a measurable first AI win — without crossing any regulatory line.

The operating principle

The lowest-scoring dimension is not where you start. The highest-leverage dimension is. Pick one win. Ship it in 90 days. Re-baseline.

You have just completed the AiBI In-Depth Readiness Assessment. Your team's report scored your institution across eight dimensions of AI readiness — strategy, governance, data, talent, culture, technology, risk posture, and use-case maturity. This playbook turns that snapshot into a 90-day execution plan.

WHAT'S INSIDE

1. Executive snapshot
2. How to read your report
3. The three 90-day paths
4. Path A · Strategy & Culture
5. Path B · Governance & Risk
6. Path C · Use Case Maturity
7. The fourth pattern · cross-cutting
8. Measurement model
9. Risk & red lines for the first win
10. Communication plan (board, exam, team)
11. The four-quarter roadmap
12. Where AiBI fits
13. Recommended next step

Who this is for

- The executive who sponsored the In-Depth Assessment.
- The compliance officer, lending lead, retail lead, or marketing lead who owns a dimension.
- The AI working group convening for the first time after the assessment.
- The board AI subcommittee asking, "OK, what now?"

What this playbook is not

- It is not your full AI strategy.
- It is not your AI policy. (Path A produces one.)
- It is not your AI use case inventory. (Path B produces one.)
- It is not a pilot vendor selection. (Path C scopes one.)

It is your **first win** — small enough to ship in 90 days, big enough to fund the next round of work, and visible enough to bring your board, your examiners, and your team along with you.

The map vs. the move

Your In-Depth report has three layers. Understanding which one to act on first is the whole game.

Layer 1 • The dimension scores

Eight scores, 0 – 100 each. This is the **map**. It tells you where the institution stands today, dimension by dimension. It does not tell you what to do.

Layer 2 • The peer band

How your institution compares to similar community banks and credit unions on each dimension. This is the **conversation starter** for the board: "We are below peer on Governance and at peer on Use-Case Maturity."

Layer 3 • The action register

Eight verbs, owners, and a 90-day window for each dimension. This is the **move**. The action register is where the playbook starts.

Read the scores once. Read the peer band before the board update. Live in the action register for 90 days.

The trap to avoid

Most institutions instinctively act on the lowest score. That feels logical and it almost always wrong. The lowest dimension is often the hardest to move (e.g., Data Readiness) and the least visible to the board. Start where a 90-day project moves you from "we should look at this" to "we shipped this and it works." That is rarely the lowest-scoring dimension.

How to pick the highest-leverage dimension

QUESTION	IF YES, THIS IS YOUR DIMENSION
Is there an executive sponsor who would champion a 90-day win here?	Strong signal
Can a visible result land in 90 days with the team you have today?	Strong signal
Will the result help fund or unblock the other dimensions?	Strong signal

QUESTION	IF YES, THIS IS YOUR DIMENSION
Is this dimension at or near peer band — meaning a small push gets us above peer?	Strong signal
Is this dimension below peer band but cheap to move?	Strong signal
Is this dimension our worst score?	Weak signal — investigate but probably wait

Pick the path that matches your top dimension

In our work with community banks and credit unions, three patterns repeat. The next three sections walk each path end-to-end. Pick the one that matches your top finding and run it. The others stay on the shelf until quarter two.

IF YOUR HIGHEST-LEVERAGE DIMENSION IS...	YOUR 90-DAY WIN IS...	SKIP TO
Strategy & Culture	An AI Acceptable Use policy signed at the board level and rolled out to every banker.	Path A
Governance & Risk	An AI use case inventory live in a working sheet, with at least 20 use cases mapped to risk tiers.	Path B
Use Case Maturity	A retail or marketing pilot in two branches / one campaign cycle, with time-saved measured.	Path C

WHAT IF TWO DIMENSIONS TIE?

Run the path with the clearer executive sponsor. Whichever leader will actually be at the day-30, day-60, and day-90 review meetings owns the win. The other dimension waits one quarter.

What every path produces

- One visible artifact (a policy, an inventory, a measured pilot).
- A documented review cadence — quarterly board update, weekly working-group meeting.
- A trained owner — the person who carries this forward in quarter two.
- An evidence packet — the file you point your examiner to when asked.
- A re-read date — the day your team re-takes the AiBI Assessment to measure delta.

PATH A · STRATEGY & CULTURE

Land an Acceptable Use policy in 90 days

For institutions where the team is ready to use AI but does not yet have permission.

OWNER: COMPLIANCE OFFICER

SPONSOR: EXECUTIVE TEAM

WIN: SIGNED POLICY + 100% ROLLOUT

Why this path

If Strategy & Culture is your highest-leverage dimension, the friction is permission, not capability. Bankers want to use AI; managers want clarity; the board wants a defensible posture. A signed Acceptable Use policy unlocks all three in one move.

Days 1 - 30 · Draft

- Convene your AI working group: a sponsor from the executive team, the compliance officer, a representative from retail and from operations.
- Start with the AiBI **Safe AI Use Checklist** and **Red / Yellow / Green Use Card**. They are the one-page anchors for the policy.
- Build a draft Acceptable Use policy. Two pages, plain English. What is allowed, what is not, who reviews.
- Identify five workflows where AI will be used in the first quarter. These become the early-adopter cases.

Days 31 - 60 · Review

- Legal & compliance review against Interagency TPRM, ECOA / Reg B, and the AIEOG AI Lexicon. For any tool that works as a model, add SR 26-2, the April 2026 model-risk guidance that superseded SR 11-7. SR 26-2 is non-binding and puts generative AI outside its formal scope, so your own risk framework governs chat tools.
- Department-head review for operational realism. If your retail leader cannot live with the policy, it will not survive contact with the branch network.
- Iterate to a board-ready version.
- Train the five early-adopter workflow owners on the approved approach.

Days 61 - 90 · Land

- Board approves the policy. Item goes in the minutes.
- Rollout: email to every employee with the one-page summary, plus a 20-minute team meeting in every department.
- First-month tracking: usage of approved tools, questions raised, exceptions logged.
- The compliance officer presents the rollout at the next quarterly board update.

WIN CONDITION

At day 90, every employee can answer the question "what does our AI policy say?" in one sentence. That is a culture win. The board sees it. Your examiner sees it. Your team feels it.

PATH B · GOVERNANCE & RISK

Stand up an AI use case inventory in 90 days

For institutions where AI is already being used — invisibly — and no one has mapped it.

OWNER: COMPLIANCE / RISK OFFICER

SPONSOR: CHIEF RISK OFFICER

WIN: 20+ INVENTORIED USE CASES

Why this path

If Governance & Risk is your highest-leverage dimension, your bank is already using AI — through vendor tools, free chat tools, and pockets of self-directed staff use — but no one has mapped it. The 90-day win is the map.

Days 1 - 30 · Discover

- Run a structured discovery interview with every department head: "where is AI showing up in your work today, in vendor tools or in informal use?"
- Audit your top 20 vendor contracts for the "AI feature" question. The Interagency TPRM Guidance applies — vendors with AI inside count.
- Build the inventory spreadsheet (use case, tool, data class, risk tier, reviewer, last reviewed).

Days 31 - 60 · Classify

- Apply the Red / Yellow / Green tiers to every use case.
- Flag any use case that touches credit decisioning, adverse action, fraud, AML, or customer authentication as Red, with documented controls (SR 26-2 principles for non-generative models).
- Add the human reviewer name and the review cadence to every Yellow row.
- Retire any use case that should not have been live in the first place — document the retirement.

Days 61 - 90 · Govern

- First quarterly AI governance update to the board. Three things: where AI lives, what changed, what is coming.
- Inventory becomes a living document — owned by the compliance officer, reviewed quarterly.
- First Yellow use case formally reviewed and signed off, creating the template for every future review.

WIN CONDITION

At day 90, your examiner can ask "where is AI used in your bank?" and you have a current, classified, reviewed answer. **55% of FIs had no AI governance framework** (Gartner via Jack Henry, 2025). That gap is your competitive moat.

PATH C · USE CASE MATURITY

Run a measurable pilot in 90 days

For institutions with permission and posture but no shipped product.

OWNER: RETAIL OR MARKETING LEAD

SPONSOR: DEPARTMENT HEAD

WIN: MEASURED TIME SAVED + QUALITY HELD

Why this path

If Use Case Maturity is your highest-leverage dimension, your bank has permission and posture but no shipped product. The 90-day win is a single working pilot in either retail or marketing, with time-saved measured.

Pick one pilot, not three

The temptation is to pilot everything. The pattern that works is one pilot, two branches (or one campaign cycle), one metric, 90 days. Use the AiBI **Retail Playbook** or **Marketing Playbook** as your pilot blueprint.

Days 1 - 30 · Set up

- Select pilot scope. Two branches for retail, or two consecutive campaign cycles for marketing.
- Approved tool in hand. Brand voice prompt (Marketing Play 1) or template library (Retail Play 2) drafted.
- Baseline measurement: hours spent on the target task today, by which roles, on what cadence.
- Compliance signs off on the workflow before live use.

Days 31 - 60 · Run

- Pilot teams use the AI assist on the target task every working day. Weekly check-in with the pilot lead.
- Track time saved per banker per week. Track quality (sample output reviews, customer feedback if relevant).
- Adjust prompts and templates weekly. The first month's results are not the last month's.

Days 61 - 90 · Decide

- Pilot review meeting with the executive team. Three artifacts: time saved, quality assessment, banker satisfaction.
- Rollout decision: full network, paused, or pivot.
- Documentation: the pilot becomes a row in your AI use case inventory, with the human reviewer named and the cadence set.

WIN CONDITION

At day 90, your bank has shipped one AI workflow that bankers like, customers cannot tell exists, and the board can point to. That is the proof point that funds Path A and Path B in the next two quarters.

Cross-cutting moves that show up in every path

Whichever 90-day path you run, four moves repeat. Plan them in week 1 of any path.

1 • Name the AI owner

Every path needs one person whose calendar reflects this work. Title varies — Compliance Officer, Head of Innovation, AI Working Group Chair — but the role is the same: this is their explicit job for 90 days.

2 • Decide your tool stance

Approved enterprise AI tool? Use it. No approved tool yet? Define the data-boundary rules so staff are not in the wrong tools on day 30.

3 • Build the evidence packet from day 1

Whichever artifact you ship, build the supporting file as you go: source documents, draft versions, reviewer sign-offs, final approval date, retention location. The examiner version is the same file.

4 • Set the re-read date

Day 90 is not the finish line. The day-180 re-take of the AiBI Assessment is where the delta becomes the next board update. Put it on the calendar in week 1.

SECTION 8 • MEASUREMENT MODEL

What to track for the 90-day win

PATH	PRIMARY METRIC	SECONDARY METRICS
A • Strategy & Culture	Board policy approved (Y/N)	% of employees briefed · exceptions logged · questions raised · usage of approved tools
B • Governance & Risk	Use cases inventoried (target: 20+)	% risk-tiered · % with named reviewer · Yellow exceptions reviewed · Red exceptions escalated
C • Use Case Maturity	Time saved per banker per week	Quality / reviewer correction rate · banker satisfaction · customer impact · rollout decision

For all three paths

- Dimension re-score at day 180 (one quarter after the win lands).
- Peer-band movement vs. the In-Depth baseline.
- Executive read-time on the quarterly board update.

Lines not to cross in the next 90 days

Each path has tempting shortcuts. Each shortcut creates real exposure. The discipline of the first win is what makes Paths 2 and 3 possible.

Path A red lines

- Do not publish an Acceptable Use policy that has not been reviewed by compliance and legal.
- Do not skip the department-head review — a policy without operational realism dies in week two.
- Do not let the policy implicitly authorize credit, AML, or adverse-action use cases. Those need separate, documented review.

Path B red lines

- Do not classify a credit-touching workflow as Yellow. Anything that bears on a credit decision is Red.
- Do not let "informal AI scoring" stay informal. If a tool ranks credit applications, treat it as a model: document it, validate it, and monitor it, consistent with SR 26-2 (which superseded SR 11-7 in April 2026) and ECOA / Reg B.
- Do not let the inventory drift to a "list we maintain quarterly." It is a living document — every new vendor contract gets an AI question.

Path C red lines

- Do not pilot a customer-facing AI workflow without documented human review.
- Do not let banker-good time savings hide a quality regression. Sample output reviews are non-negotiable.
- Do not roll out the pilot before the day-90 review meeting. The discipline of a hard checkpoint is what produces a real decision.

The 90-day win is small. Its credibility is large. Anything that compromises the credibility kills the next three quarters of investment.

Bring the board, the exam, and the team along

Board update — three slides, every quarter

1. **Where AI is used today.** Total use cases (Path B), or policy adoption metrics (Path A), or pilot results (Path C).
2. **What changed since last quarter.** New use cases, exceptions, incidents, learnings.
3. **What is coming.** Next dimension under work, next 90-day win, training needs, vendor reviews.

Examiner readiness — three answers

1. **"What does your AI policy say?"** One sentence. (Path A produces this.)
2. **"Where is AI used in your bank?"** Pointer to the inventory. (Path B produces this.)
3. **"Who reviews the output?"** Pointer to the human reviewer in the inventory. (All three paths produce this.)

Team rollout — one cadence

- Day 30: AI working group meets weekly. Department heads informed.
- Day 60: Cross-functional review with compliance, legal, IT/InfoSec, retail, marketing, lending.
- Day 90: Rollout briefing — every employee gets the one-pager, every department gets a 20-minute meeting.
- Day 180: Re-baseline against the AiBI Assessment. Communicate the delta to the board and the team.

From first win to AI-capable bank

The 90-day win is the start of a different conversation, not the end of one. Once your bank has a documented policy, a live use case inventory, or a measurable pilot, the next four quarters compound.

QUARTER	OUTCOME	WHAT'S NEW ON THE CALENDAR
Q1 — your first win	Path A, B, or C complete. Anonymized team rollup of assessment results in the executive team's hands.	Quarterly board update established. Working group cadence set.
Q2 — second win	Whichever path was not Q1 gets the next 90-day plan. Foundations course rolled out to your 20 highest-leverage bankers.	Foundations cohort #1. Day-180 re-read of the AiBI Assessment.
Q3 — third win	The remaining path. First AiBI Specialist (AiBI-S) certificants in compliance, retail, or marketing.	AiBI-S enrollment for role leads. First measured peer-band movement.
Q4 — re-baseline	Re-take the AiBI Readiness Assessment as an institution. Score change becomes a published metric for the board.	Year-over-year delta in the annual report.

What the four quarters give you

- A signed AI policy + an AI use case inventory + at least one shipped AI workflow.
- A cohort of 20+ trained bankers (Foundations) and 3 - 5 role specialists (AiBI-S).
- A measurable institutional movement on the AiBI Assessment — the score becomes the board's KPI.
- An examiner-ready file: policy, inventory, workflow SOP, human review evidence, board cadence.

The institute's role in the four-quarter arc

AIBI OFFERING	WHERE IT FITS	WHO IT'S FOR
Free Readiness Assessment	Day 0 + Day 180 + Day 365 — the score is the institutional KPI.	Executive sponsor, board, department heads.
In-Depth Readiness Assessment (\$99)	Day 0 — the deeper instrument that produced this playbook.	The In-Depth subject and their working group.
AiBI Foundations course (\$295/seat)	Quarters 2 – 3 — the course that turns curious bankers into capable ones.	The 20 – 50 bankers who will carry AI work forward.
AiBI Specialist (AiBI-S)	Quarter 3+ — role-specific deepening for compliance, lending, retail, ops, marketing leaders.	3 – 5 role leads per institution.
AiBI Leadership Advisory	Quarters 1 – 4 — fractional Chief AI Officer engagement.	Institutions running cohorts at scale.

What's deliberately not included

- Vendor selection. We don't sell AI tools. We help you decide which approved tools fit your workflow.
- Custom AI build. We don't build AI products. We build the bankers who run them.
- Replacement for compliance, legal, or risk. We support those functions. We don't substitute for them.

SECTION 13 • RECOMMENDED NEXT STEP

What to do in the next two weeks

1. Share this playbook with your executive sponsor and your AI working group.
2. Pick your path. One of A, B, or C — not all three.
3. Name the owner. The person whose calendar reflects this work for 90 days.
4. Book the day-30, day-60, and day-90 review meetings now, before the work starts.
5. Book an Executive Briefing at **AlBankingInstitute.com**. Bring this playbook. We will walk through your dimension scores, your chosen path, and the support that fits your next four quarters.

**A reading is a snapshot. The value of the snapshot is the next snapshot.
Set the re-read date now. The delta is what you brief leadership on.**

Sources: AiBI In-Depth Readiness Assessment (v2 instrument); GAO-25-107197 (May 2025); AIEOG AI Lexicon (US Treasury / FBIIC / FSSCC, Feb 2026); SR 26-2, Revised Guidance on Model Risk Management (FRB / OCC / FDIC, April 17, 2026), which superseded SR 11-7 (2011); Interagency Third-Party Risk Management Guidance; ECOA / Regulation B; FDIC Quarterly Banking Profile Q4 2024; Gartner Peer Community data via Jack Henry, "Getting Started in AI" (2025).

This playbook is a starter artifact and planning aid, not legal advice or an examiner-approved roadmap. Adapt before adoption and validate final sequencing against your institution's risk appetite, budget, vendors, and board priorities.