

ROLE CONTROL KIT - LENDING - V2.1

The Lending AI Control Kit

Five controls for using AI in lending without losing explainability, file support, or human accountability.

The operating principle

AI prepares the workbench. The human underwriter owns the decision. The file proves the reason.

AI can extract, summarize, compare, calculate, draft, and organize. It cannot become the unreviewed decision-maker in a credit workflow.

The Five Lending AI Controls

#	CONTROL	EVIDENCE TO RETAIN
1	No AI credit decisioning without formal review.	Approved use case, model/tool review, decision authority, control owner.
2	Every principal reason traces to the file.	Principal reason table, source document, reviewer sign-off.
3	AI drafts only from human-provided facts.	Prompt, approved input set, AI output, reviewer edits.
4	Every customer-facing or credit-impacting output has human review.	Review checklist, final approval, escalation log.
5	Every pilot has evidence, metrics, and stop trigger.	Pilot scorecard, exception log, scale / revise / stop decision.

WHAT'S INSIDE

- | | |
|---|--|
| 1. Executive snapshot | 10. Model risk & continuous monitoring |
| 2. The Five Lending AI Controls | 11. Human-in-the-loop & fallbacks |
| 3. Why lending is different | 12. Five lending plays |
| 4. The governed lending operating model | 13. Editable templates |
| 5. Prepare vs. decide matrix | 14. 30-day lending pilot plan |
| 6. Modern AI lending architecture | 15. Measurement model |
| 7. Alternative data guardrails | 16. What good looks like |
| 8. Explainability & adverse action | 17. 45-minute control sprint |
| 9. Fair lending governance controls | 18. Final operating principle |

Who this is for

- Chief Lending Officers, Credit Officers, Loan Operations leaders.
- Underwriters.
- Fair Lending / CRA / Compliance teams.
- Model risk management partners.
- Credit administration teams.
- AI governance committees.
- Executives evaluating AI-assisted lending workflows.

The core problem

Lending is the highest-stakes AI lane in community banking. A weak AI workflow in marketing may create awkward copy. A weak AI workflow in lending can create regulatory exposure, borrower harm, fair-lending risk, adverse-action failures, model risk issues, and examiner scrutiny.

The opportunity is real. Lending workflows are full of repetitive, document-heavy, language-heavy, calculation-heavy work: collecting and summarizing financial documents, organizing credit memos, reviewing loan files for missing information, drafting internal narratives, preparing borrower-friendly explanations, comparing policy requirements, checking fair-lending language, and building decision packet indexes. But the control requirement is stricter than almost anywhere else.

The opportunity

A safe AI lending program can help teams reduce document-review friction; improve credit memo consistency; create clearer decline explanations; standardize decision packet documentation; catch weak or generic adverse-action language; identify fair-lending language concerns before finalization; improve coaching and underwriting quality; and reduce rework between lending, credit, and compliance.

Where AI mistakes become regulator letters

The Equal Credit Opportunity Act and Regulation B require that adverse-action notices provide specific, accurate principal reasons for the action. The reasons must be grounded in the file, not invented by a model and not hidden inside a black-box recommendation.

The black-box excuse is over

If a creditor uses a complex algorithm, AI system, or machine learning model in a credit decision, the institution still has to provide the applicant with specific reasons for adverse action.

If the bank cannot explain the reason, the bank cannot safely use the system to support the decision.

RED LINE

No AI system writes, sends, or finalizes an adverse-action notice without documented human review. No AI system independently approves, denies, ranks, prices, or recommends credit action without model risk, fair lending, compliance, and human decision controls. Anything that bears on a credit decision must be traceable to the file.

The lending workbench, not the lending oracle

A modern AI lending workflow should not be one giant model asked to "underwrite the borrower." That is the wrong architecture. A safer pattern is a **governed lending workbench**: specialized AI components assist with bounded tasks, and a human underwriter remains accountable for the final decision.

Six layers

- **Layer 1 · Data boundary.** Define exactly what data can be used, where, and by whom (public product info, internal lending policy, redacted file notes, customer NPI/PII, credit-bureau data, income/employment, tax returns & bank statements, protected-class proxy-risk data, final credit decision records).
- **Layer 2 · Task-specific AI assistance.** Extracting fields, summarizing redacted notes, organizing credit memo sections, drafting borrower-friendly explanations, checking language clarity, highlighting missing information, preparing index of file evidence.
- **Layer 3 · Policy and eligibility checks.** AI may compare file facts against policy, but the system must not silently decide eligibility. Checks must be traceable, source-linked, reviewed, flagged when

uncertain, and retained.

- **Layer 4 · Human decision gate.** A qualified human underwriter, credit officer, or authorized decision-maker makes the final approval, denial, counteroffer, or condition decision.
- **Layer 5 · Explainability and evidence.** Every adverse decision must be supported by specific principal reasons grounded in the file. Retain file source, principal reasons, reviewer sign-off, AI draft if used, human edits, final adverse-action language, fair-lending review result.
- **Layer 6 · Monitoring and validation.** If AI or models are used in credit-related workflows, monitoring must be risk-based and ongoing: performance stability, overrides, drift indicators, adverse-action reason patterns, protected-basis outcome gaps where applicable, exception rates, data quality, vendor changes.

What AI can prepare and what humans must decide

WORKFLOW	AI CAN PREPARE	HUMAN MUST DECIDE	EVIDENCE RETAINED
Document review	Extract fields, summarize missing items	Whether file is complete	Source docs, extraction output, reviewer
Credit memo	Draft structure, summarize facts, organize exhibits	Final credit narrative	Draft, reviewer edits, final memo
Policy comparison	Flag where file appears to meet or miss policy	Whether exception or approval is warranted	Policy source, file facts, reviewer decision
Adverse action	Draft language from human-provided principal reasons	Final reasons and final notice	Principal reasons, source file, reviewer, final notice
Fair-lending pre-check	Flag potentially problematic phrases or proxy-risk language	Whether language or process creates fair-lending concern	Output, reviewer, final disposition
Risk calculations	Calculate ratios or summarize model output	Whether model result is relied on	Data source, calculation, validation status
Decision packet	Create index and assemble references	Whether packet is accurate and complete	Index, source references, reviewer sign-off
Borrower summary	Draft plain-English explanation after decision	Final communication and relationship handling	Approved letter, call script, reviewer

Multi-agent lending workbench

A lending workflow may use specialized AI agents or services, each with a narrow job. The safest pattern is **not** "one AI underwriter." The safer pattern is: one agent extracts, one agent checks policy, one agent summarizes risk, one agent drafts the memo, one orchestrator organizes the packet, one human underwriter decides.

DOCUMENT AGENT

Extracts and structures information from application PDFs, tax returns, bank statements, financial statements, or borrower-provided documents.

Controls: approved environment only · source document references · confidence flag · human verification of extracted values.

POLICY AGENT

Compares file information against lending policy.

Controls: cites policy section · flags uncertainty · cannot approve exceptions · cannot make final eligibility decision.

FRAUD / ANOMALY AGENT

Flags inconsistencies, missing information, mismatched names, unusual document patterns, or potential fraud indicators.

Controls: cannot accuse borrower · cannot close investigation · routes issues to human review.

Specialized agents · risk and orchestration

RISK CALCULATION AGENT

Computes or summarizes ratios and model outputs.

Controls: source data traceability · calculation review · model validation status · no unreviewed decisioning.

ORCHESTRATOR AGENT

Assembles the work product into a draft credit proposal or decision packet.

Controls: cannot approve, deny, or send notices · must show source references · routes to human underwriter · logs evidence.

Useful, high-risk, and the in-between

Alternative data can expand access to credit, but it can also create fairness, privacy, explainability, and proxy-discrimination risk.

Potentially useful when properly governed

- Cash-flow data.
- Rent payment history.
- Utility payment history.
- Verified income trends.
- Deposit account behavior.
- Business revenue consistency.

High-risk or generally problematic

- Social media profiles.
- Education background.
- Location or neighborhood proxies.
- Browsing behavior.
- Device metadata.
- Non-transparent third-party scores.
- Variables that may proxy protected-class characteristics.

Alternative data review checklist

- ❑ The data has independent predictive value.
- ❑ The data is permitted by policy and law.
- ❑ The data source is reliable and documented.
- ❑ If this data contributes to denial, pricing, terms, or conditions, can we explain it to the applicant and trace it to a file source?
- ❑ The data is reviewed for protected-basis outcome gaps, proxy-variable risk, explainability, privacy, and state-law concerns where applicable.
- ❑ The data is subject to model-risk review if used in a decisioning model.
- ❑ Human review and override are documented.

Required adverse-action workflow

AI may not select, infer, rank, or invent principal reasons. AI may only help express principal reasons already identified by a qualified human reviewer and supported by the file.

1. Human underwriter identifies the actual principal reasons from the file.

2. AI may draft clearer wording using only those reasons.
3. Human reviewer compares every AI-generated sentence to the source file.
4. Unsupported, vague, or invented language is removed.
5. Final notice uses approved language and accurate principal reasons.
6. Evidence is retained.

PROHIBITED PATTERN

"Why should this borrower be denied?"

APPROVED PATTERN

"Using only the underwriter-provided principal reasons below, draft clear adverse-action language. Do not add reasons. Do not infer facts. Flag any reason that is vague or unsupported."

What can go wrong

- Historical bias is learned by the model.
- Neutral variables act as proxies for protected-class characteristics.
- Alternative data creates digital redlining risk.
- Model outputs are accepted without human challenge.
- Adverse-action reasons become vague or generic.
- Protected-class impact is not monitored.

Review approach

For AI-assisted lending workflows, monitor for protected-basis outcome gaps, proxy-variable risk, adverse-action reason patterns, override patterns, and explainability issues as part of fair-lending governance, reputation risk, state-law, Fair Housing Act, mortgage, and institution-specific control review where applicable.

If a variable cannot be explained, defended, tested, and translated into a specific principal reason when relevant, it should not drive an adverse action.

What to monitor

Signals to track

- Population stability.
- Input data quality.
- Model performance.
- Override frequency.
- Exception rates.
- Adverse-action reason distribution.
- Approval / denial pattern changes.
- Drift indicators.
- Vendor model updates.
- Data source changes.
- Complaint trends.
- Fair-lending review signals.

Possible monitoring tools

- Population Stability Index (PSI).
- Challenger models.
- Champion / challenger comparison.
- Back-testing.
- Benchmark review.
- Override testing.
- Sample file review.
- Outcome monitoring.
- Adverse-action reason audits.

Important qualification

Current model-risk guidance, including SR 26-2 where applicable, should be considered when AI supports scoring, pricing, risk ranking, eligibility, fraud, or other model-like decision processes. Generative AI used for drafting, summarization, file organization, or language support should still be governed through approved-tool controls, data boundaries, human review, evidence retention, vendor oversight, and incident response.

Where AI stops

Human-in-the-loop means

A qualified human understands the workflow, has authority to approve or reject the output, compares the AI output to source evidence, can override the AI, signs off on final use, and is accountable for the decision.

Required human gates — review before

- Loan approval.
- Loan denial.
- Counteroffer.
- Adverse-action notice.
- Credit memo finalization.
- Exception approval.
- Pricing decision.
- Condition removal.
- Customer-facing explanation.
- Regulatory or examiner response.

Deterministic fallback — route to a human when

- Confidence is below threshold.
- Source document is missing.
- Model output conflicts with file evidence.
- Protected-class proxy risk is flagged.
- Adverse-action reason is vague.
- Fraud signal appears.
- Policy exception is needed.
- Customer-facing language is generated.
- The AI cannot cite a source.

PLAY 1

Adverse-action language tuning

Make adverse-action language clearer without losing file-specific principal reasons.

Why it matters

Underwriters often identify the right reasons but struggle to translate them into clear, borrower-facing language. AI can help polish the language. It cannot create the reason.

Approved inputs

- Human-provided principal reasons.
- Approved adverse-action template.
- Redacted file references.
- Product type.
- Reviewer instructions.

Prohibited inputs

- Unredacted customer identifiers in public tools.
- Full credit file in public tools.
- Protected-class information.
- Unsupported reasons.
- Model output with no file source.

PROMPT TEMPLATE · ADVERSE-ACTION LANGUAGE

You are assisting a human underwriter with draft adverse-action language.

Use only the principal reasons provided below. Do not add reasons. Do not infer facts. Do not make the decision. If a reason is vague or unsupported, flag it as [VERIFY].

Output:

1. Clear borrower-facing language
2. Source reason used
3. Any unsupported or vague reason flagged [VERIFY]
4. Reviewer checklist

Principal reasons from the file:
 [INSERT HUMAN-PROVIDED REASONS]

Approved template language:
[INSERT TEMPLATE]

Output: draft language · reason traceability · verify flags · reviewer checklist. **Review:** underwriter and compliance reviewer compare the language to the file. **Evidence:** principal reasons · draft AI language · human edits · final adverse-action notice · reviewer sign-off. **Metrics:** minutes saved per declined application · unsupported AI phrases removed · reviewer correction rate · adverse-action exception findings.

PLAY 2

Borrower-friendly decline summary

Plain-English borrower conversation summary after adverse-action language has been approved.

Why it matters

The regulatory notice satisfies the legal requirement. The borrower conversation affects the relationship, complaint risk, and future application quality.

Approved inputs

- Approved adverse-action reasons.
- Product type.
- General next-step guidance.
- Relationship banker tone guidance.

Prohibited inputs

- New reasons not in the notice.
- Promises of future approval.
- Unapproved credit counseling.
- Sensitive file detail beyond the approved explanation.

PROMPT TEMPLATE · BORROWER DECLINE CONVERSATION

Using only the approved adverse-action reasons below, draft a plain-English conversation guide for the loan officer.

Do not add reasons. Do not promise approval. Do not provide legal or credit counseling. Keep the tone respectful and specific.

Output:

1. Short opening
2. Plain-English explanation
3. What the borrower can ask about
4. What the loan officer should not say
5. Escalation note

Approved reasons:

[INSERT APPROVED REASONS]

Output: call guide · do / do-not-say list · escalation note. **Review:** loan officer and compliance approve before use. **Evidence:** approved notice · conversation script · reviewer approval. **Metrics:** complaint rate · borrower follow-up quality · loan officer confidence · call prep time saved.

PLAY 3

Fair-lending pre-check

AI as a second reader to flag potentially problematic language before the packet is final.

Why it matters

Fair-lending issues are often easier to prevent than fix after the fact.

Approved inputs

- Redacted underwriting notes.
- Draft adverse-action language.
- Approved policy excerpts.
- Phrase lists.

Prohibited inputs

- Protected-class information.
- Unredacted customer file data in public tools.
- Unsupported model explanations.
- Final decision without review.

PROMPT TEMPLATE · FAIR-LENDING LANGUAGE SCREEN

You are reviewing draft lending language for potential fair-lending concerns.

Identify language that could be vague, generic, unsupported by file facts, or potentially interpreted as proxying for a protected-class concern.

Do not make a legal determination. Do not certify compliance.

Output:

1. Phrase to review
2. Why it may be concerning
3. Suggested neutral alternative
4. Whether file support is needed
5. Escalation recommendation

Draft language:

[INSERT REDACTED DRAFT]

Output: phrase screen · neutral wording suggestions · escalation flags. **Review:** compliance owns the final interpretation. **Evidence:** draft language · AI screen · compliance decision · final language. **Metrics:** problematic phrases flagged · compliance edits per packet · first-pass approval rate · training opportunities identified.

PLAY 4

Examiner-ready decision packet index

Organize the credit decision packet into a one-page index that points to source evidence.

Why it matters

Examiners are not impressed by document volume. They are impressed by clarity, traceability, and source evidence.

PROMPT TEMPLATE · DECISION PACKET INDEX

Create a one-page decision packet index using only the source list and summaries below.

Do not create new evidence. Do not make a credit decision. Point to source documents that already exist.

Output:

1. Decision summary
2. Source document index
3. Principal reasons and file support
4. Reviewers and sign-off
5. Open items
6. Retention note

Source list:

[INSERT SOURCE LIST]

Decision summary:

[INSERT HUMAN-APPROVED SUMMARY]

Output: decision packet index · source map · open item list. **Review:** credit administration and compliance. **Evidence:** source list · draft index · final index · reviewer sign-off. **Metrics:** decision packet review time · missing document exceptions · examiner / audit follow-up requests · rework reduction.

PLAY 5

Lending language coaching guide

One-page guide of language that holds up vs. language that does not.

Why it matters

The same language issues appear repeatedly across loan officers, underwriters, and branches. AI can help cluster examples, but lending leadership owns the final coaching standard.

PROMPT TEMPLATE · LENDING LANGUAGE COACHING

Review these anonymized lending language examples.

Cluster them into:

1. Clear and file-specific
2. Too generic
3. Potential proxy concern
4. Missing source support
5. Better alternative language

Do not make a legal determination. The lending lead and compliance reviewer will curate final examples.

Examples:

[INSERT ANONYMIZED LANGUAGE]

Output: coaching guide draft · phrase clusters · better-language examples · review flags. **Review:** lending lead and compliance approve. **Evidence:** anonymized examples · AI cluster output · curated final guide · approval date. **Metrics:** coaching guide adoption · fewer repeated language issues · faster new-lender onboarding · compliance review improvement.

Eight reusable lending artifacts

Template 1 · AI lending use-case intake form

FIELD	RESPONSE
Use case name	
Lending workflow affected	
Business purpose	
Proposed tool	
User role	
Input data class	
Output type	
Does it influence approval, denial, pricing, or terms?	Yes / No
Customer-facing output?	Yes / No
Adverse action involved?	Yes / No
If denial, pricing, terms, or conditions are affected, can each factor be explained and traced to the file?	Yes / No / TBD
Model risk review required?	Yes / No / TBD
Fair-lending review required?	Yes / No / TBD
Human reviewer	
Approval checkpoint	
Retention location	
Risk tier	Green / Yellow / Red
Decision	Approved / Conditional / Rejected

Template 2 · Principal reason traceability table

PRINCIPAL REASON	FILE SOURCE	AI DRAFT SENTENCE	HUMAN REVIEWER EDIT	FINAL LANGUAGE
------------------	-------------	-------------------	---------------------	----------------

Template 3 · Adverse-action AI review checklist

CHECKPOINT	OWNER	STATUS	NOTES
Principal reasons were identified by a qualified human reviewer.			
Each reason has file support.			
AI did not add, rank, infer, or remove reasons.			
Customer-facing language matches approved template standards.			
Final notice and evidence packet are retained.			

Template 4 · Alternative data review worksheet

FIELD	RESPONSE
Data source and owner	
Decision element affected	Denial / Pricing / Terms / Conditions / Other
File source and applicant explanation	
Protected-basis / proxy-variable review completed	Yes / No / TBD
FCRA, FHA, GLBA, and state-law checkpoints considered where applicable	Yes / No / TBD
Human override and exception path	

Template 5 · Fair-lending phrase screen

PHRASE	CONCERN	SOURCE SUPPORT NEEDED?	SUGGESTED NEUTRAL LANGUAGE	REVIEWER DECISION
--------	---------	------------------------	----------------------------	-------------------

Template 6 · Decision packet index

PACKET SECTION	SOURCE DOCUMENT	OWNER	STATUS	NOTES
Application snapshot				
Underwriting summary				
Principal reasons				
Adverse-action notice				
Fair-lending review				
Reviewer sign-off				

Template 7 · Human review sign-off log

DATE	WORKFLOW	REVIEWER	DECISION AUTHORITY	APPROVE / REVISE / REJECT	RETENTION LOCATION
------	----------	----------	-----------------------	------------------------------	-----------------------

Template 8 · 30-day lending AI pilot scorecard

METRIC	BASELINE	TARGET	ACTUAL	SCALE / REVISE / STOP NOTE
Files tested				
Minutes saved per file				
Reviewer correction rate				
Unsupported AI phrases removed				
Traceability completeness				
Stop trigger reached?	No	No		

Lending AI evidence packet

For each AI-supported lending workflow, retain:

- Use-case intake form.
- Tool approval status.
- Data classification.
- Prompt or instruction.
- Source material description.
- AI output.
- Human reviewer.
- Reviewer changes.
- Principal reason traceability table, if adverse action is involved.
- Final approved output.
- Approval date.
- Retention location.
- Next review date.

Select, run, screen, decide

Week 1 — select and control the pilot

Outcomes: approved tool selected; lending and compliance agree on red lines; Safe AI Use Checklist distributed to lending team; first use case selected; human reviewer named.

Recommended first use case: adverse-action language tuning with human-provided principal reasons — or decision packet index creation.

Artifacts: AI lending use-case intake form · approved prompt · review checklist.

Week 2 — run five controlled files

Outcomes: pilot run on five declined applications or five redacted sample files; underwriter time logged; reviewer corrections tracked; unsupported AI language removed.

Artifacts: principal reason traceability table · AI draft output · reviewer edits.

Week 3 — add fair-lending screen and packet index

Outcomes: fair-lending phrase screen added; first decision packet index produced; compliance reviews pilot outputs.

Artifacts: fair-lending phrase screen · decision packet index · evidence packet.

Week 4 — decide whether to scale

Outcomes: lending lead and compliance review pilot; time saved and quality reviewed; go / revise / stop decision made; training needs identified.

Artifacts: 30-day lending AI pilot summary · approved workflow SOP · next-step plan.

Productivity, quality, risk, adoption

Productivity

- Minutes saved per file.
- Minutes saved per decline.
- Credit memo drafting time.
- Decision packet assembly time.
- Review turnaround time.

Quality

- Reviewer correction rate.
- Unsupported AI phrases removed.
- Missing document exceptions.
- Fair-lending flags.
- First-pass compliance approval rate.

Risk

- Adverse-action exceptions.
- Model / tool incidents.
- Override rate.
- Customer complaints.
- Protected-basis outcome gap signals.
- Source traceability completeness.

Adoption

- Underwriters trained.
- Prompts reused.
- Decision packets indexed.
- Coaching guide adoption.
- Weekly usage of approved workflow.

Ten questions the lending team can answer at day 30

1. Which lending AI use cases are allowed?
2. Which are prohibited?
3. Which tool is approved?
4. What data may be entered?
5. Who reviews AI output?
6. How are principal reasons traced to the file?
7. What happens when AI output is unsupported?
8. Where is the final evidence packet retained?
9. What metric proves value?
10. What must be reviewed before scale?

The goal is not AI underwriting. The goal is AI-assisted lending operations that are faster, clearer, more consistent, explainable, reviewed, traceable, fair-lending aware, and file-supported.

SECTION 16 · 45-MINUTE CONTROL SPRINT

Next step: Run the 45-Minute Lending AI Control Sprint

Pick one lending workflow, classify the data, define the human gate, map principal reasons to the file, build the evidence packet, and decide whether the pilot is safe to launch.

1. Take the AI Readiness Assessment.
2. Assign lending, credit, and compliance leads to the Foundation Course.
3. Complete the Red / Yellow / Green Use Card and Safe AI Use Checklist.
4. Build the Lending AI Use-Case Intake Form.
5. Run the adverse-action language sandbox using fictional or redacted data.
6. Build the Principal Reason Traceability Table.
7. Save the lending workflow SOP to the Toolbox.
8. Reassess after 90 days.

Faster lending, never less explainable

AI can make lending faster. But lending cannot become less explainable. The lending team still owns the decision. Compliance still owns the review. The file still owns the reason.

AI prepares. Humans decide. The file proves.

Sources: ECOA / Regulation B; CFPB Circular 2022-03 on adverse action and complex algorithms; FCRA; Fair Housing Act where mortgage or housing-related lending is involved; GLBA and customer information safeguards; SR 26-2 where applicable; Interagency Third-Party Risk Management Guidance; Financial Services AI Risk Management Framework; AIEOG AI Lexicon (US Treasury / FBIIC / FSSCC, February 2026); GAO-25-107197 (May 2025); AiBI Foundations curriculum.

This kit is a starter artifact, not legal advice or an examiner-approved lending procedure. Adapt before adoption and route final lending, adverse-action, and fair-lending language through compliance, credit, legal, and risk review.