



The AI Banking Institute
Turning Bankers into Builders

CHEAT SHEET · V2.0

The Banker's Prompt Strategy Cheat Sheet

Better prompts. Safer inputs. More reusable outputs. A practical prompting reference for banking professionals who need useful AI output without leaking data, weakening review discipline, or turning a chat response into an unapproved business record.

One question before every prompt

Strip the specifics. Keep the structure. Ask clearly. Verify before use.

Before writing any prompt, ask one question: **Would I be comfortable if this prompt appeared outside the bank?** If the answer is no, stop and either strip the sensitive details, use an approved internal tool, or escalate.

WHAT'S INSIDE

1. The 10-second safety spine
2. Green / yellow / red prompt safety
3. The basic prompt recipe
4. The RCFC + R banker pattern
5. The enterprise prompt standard
6. The trust-check reflex
7. Daily prompt cards for bankers
8. The style and context brief
9. Advanced: prompt chaining
10. Advanced: the XML sandwich
11. Advanced: rationale, not hidden reasoning
12. Advanced: few-shot examples
13. Advanced: assistant prompt anatomy
14. The prompt escalator
15. Prompt quality checklist
16. Prompt-to-skill upgrade
17. The final rule

The core safety move

INSTEAD OF THIS

John Smith has account 123456 and a \$42,000 CD maturing on June 14.

USE THIS

[CUSTOMER] has [ACCOUNT TYPE] with [AMOUNT] maturing on [DATE].

Safe placeholders

- [CUSTOMER], [FIRST NAME]
- [ACCOUNT TYPE], [ACCOUNT NUMBER]
- [DATE], [AMOUNT]
- [BRANCH], [PRODUCT]
- [DOCUMENT], [POLICY SECTION]
- [NEXT STEP], [REVIEWER]

Never put these in a public AI tool

- Customer names · account numbers · SSNs · DOBs.
- Addresses, phone numbers, email addresses.
- Account balances · transaction details · credit file details.
- SAR / AML investigation information.
- Authentication details.
- Confidential examiner correspondence.
- Legal advice requests.
- Confidential vendor contracts.
- Passwords, API keys, tokens, system credentials.

Three lanes for every prompt

Green — usually safe

Use when the prompt contains no sensitive data and the output is not customer-impacting. Examples: draft a fictional training scenario; summarize a public regulator bulletin; brainstorm branch huddle topics; rewrite a generic internal note; create a blank checklist; build a prompt template.

Rule: if you could write it on a whiteboard in the break room, it is usually Green.

Yellow — approved tool and human review required

Use when the output could be seen by a customer, examiner, manager, or downstream business process. Examples: customer-facing email copy; summarize internal procedures; rewrite campaign language; create operational checklists; prepare meeting summaries from redacted notes; draft policy training materials.

Rule: if a customer, examiner, or decision-maker could read it, a qualified human reviews it before use.

Red — do not use public AI tools

Use an approved enterprise workflow only, if allowed at all. Examples: credit decisions; adverse-action reasoning; SAR narratives; AML alert disposition; fraud conclusions; authentication decisions; legal or compliance determinations; customer-specific financial advice; account-specific analysis.

Rule: if the output could affect a customer, financial decision, control, or regulatory filing, treat it as Red until Compliance says otherwise.

Four moves to get from generic to useful

1 · Say what you want

WEAK

Help with this.

BETTER

Rewrite this internal email so the action, owner, and deadline are clear.

2 · Give the background

Tell the AI who the audience is, what the situation is, and what matters. Example: "This is for branch managers. The tone should be warm, direct, and operational. The goal is to reduce confusion before Monday morning."

3 · Show what good looks like

Give the format, example, or success criteria. Example: "Output three bullets, each under 20 words, followed by one closing sentence."

4 · Refine what comes back

Do not just say "try again." Improve the prompt. Better refinements:

- Add a clearer role.
- Add missing context.
- Tighten the format.
- Add constraints.
- Ask for options.
- Add review criteria.

SECTION 4 · THE RCFC + R BANKER PATTERN

Five lines, every time

LINE	STANDS FOR	WHAT GOES HERE
R	Role	Who should the AI act as?
C	Context	What background does it need?

LINE	STANDS FOR	WHAT GOES HERE
F	Format	What should the output look like?
C	Constraints	What must it not do?
R	Review	Who reviews it before use?

If your prompt does not start with "You are...," you are probably leaving quality on the table.

WEAK PROMPT

Write a follow-up email.

STRONG BANKING PROMPT

Role: You are a relationship banker at a community bank. Context: [CUSTOMER] missed a follow-up call about [PRODUCT]. Use placeholders only. Format: One short email under 120 words with a clear next step. Constraints: Do not invent rates, fees, deadlines, eligibility, or promises. No exclamation points. Review: The banker will personalize and send. This is a draft.

Three required elements for higher-risk work

1 · Situational context

"A branch manager needs to explain a policy update to frontline staff before tomorrow's huddle."

2 · Professional role

"You are a retail banking operations trainer writing for frontline branch staff."

3 · Explicit constraints

"Do not change policy meaning. Do not add requirements. Flag unclear sections with [VERIFY]. Include escalation triggers."

SECTION 6 · THE TRUST-CHECK REFLEX

Verify four things before using AI output

1 · Numbers

Rates, fees, APYs, balances, dates, percentages, dollar amounts, counts, ratios.

If a number matters, verify it at the source.

2 · Dates

Effective dates, deadlines, maturity dates, statement dates, policy dates, filing dates, review dates.

AI can be confidently wrong about time.

3 · Names and citations

Regulation names, agency names, document titles, policy names, source citations, quoted language.

If AI cites it, look it up.

4 · Policy & compliance claims

Internal policy requirements, legal standards, compliance obligations, product disclosures, adverse-action language, BSA/AML statements, privacy requirements.

AI can summarize. It cannot be the final compliance reviewer.

Six reusable cards

CARD 1

Rewrite Card

Use when an internal message is unclear.

PROMPT

You are a banking operations editor.

Rewrite the message below so the action, owner, deadline, and escalation path are clear. Keep the meaning the same. Do not add new policy requirements. Flag unclear items with [VERIFY].

Output:

1. Revised message
2. Action owner
3. Deadline
4. Open questions
5. Escalation note

Message: [PASTE SAFE / APPROVED TEXT]

CARD 2

Summarize Card

Use for meetings, huddles, and internal notes.

PROMPT

You are helping a bank team create a structured meeting summary.

Use only the notes below. Do not add facts. Do not include customer identifiers. Flag missing owners or dates with [VERIFY].

Output:

1. Key decisions
2. Action items
3. Owners
4. Due dates
5. Open questions
6. Risks or escalation items

Notes: [PASTE REDACTED NOTES]

CARD 3

Document Translation Brief

Use when dense procedures need to become role-specific job aids.

PROMPT

You are a banking procedure translator writing for [ROLE].

Translate the procedure section below into plain language. Do not change policy meaning. Do not add requirements. Flag unclear or missing information with [VERIFY].

Output:

1. What changed
2. What [ROLE] needs to do
3. What not to do
4. When to escalate
5. Manager checklist

Procedure section: [PASTE APPROVED PROCEDURE TEXT]

CARD 4

Review Checklist Card

Use when AI-generated output needs review.

PROMPT

You are a banking risk reviewer.

Create a review checklist for the AI-generated output below. Check for accuracy, data sensitivity, customer impact, disclosure needs, compliance risk, unsupported claims, and retention requirements.

Output:

1. Review checklist
2. Risk flags
3. Required source checks
4. Approval owner
5. Retention note

Output to review: [PASTE AI OUTPUT]

CARD 5

Options Card

Use when you need choices, not one answer.

PROMPT

You are helping a banking team compare options.

Create three options for [TASK]. For each option, include:

1. Best use case
2. Strength
3. Risk
4. Review need
5. Recommendation

Constraints: Do not invent data. Flag missing context with [VERIFY].

Task: [INSERT TASK]

CARD 6

Manager Coaching Card

Use when a manager needs coaching language.

PROMPT

You are a branch manager preparing coaching notes for a frontline banker.

Use the fictional scenario below. Do not include real customer data. Keep the tone supportive and specific.

Output:

1. What the banker did well
2. Coaching opportunity
3. Suggested phrasing
4. Practice scenario
5. Escalation reminder

Scenario: [INSERT FICTIONAL SCENARIO]

If you repeat context, save it as a brief

If you repeat the same background more than three times a week, create a reusable brief. Paste it at the top of a session or save it as a custom assistant instruction.

BANKER STYLE BRIEF TEMPLATE

I work as [ROLE] at a [BANK / CREDIT UNION TYPE].

My audience is usually [AUDIENCE].

My writing style should be:

- [TRAIT 1]
- [TRAIT 2]
- [TRAIT 3]

Avoid:

- hype
- unsupported claims
- legal conclusions
- compliance determinations
- invented rates, fees, dates, or deadlines
- customer-specific facts

Always:

- ask for missing context when needed
- use placeholders instead of customer data
- flag uncertain items with [VERIFY]
- identify review owner when output could be used at work
- label output as draft unless approved

My preferred output format is:

[INSERT FORMAT]

Approved examples of good output:

[INSERT EXAMPLES]

Split complex banking work into stages

Do not ask one prompt to do everything. For complex banking work, split the task into smaller prompts.

Example · procedure-to-job-aid chain

PROMPT 1 · EXTRACT

Identify the key steps, roles, exceptions, and terms from this procedure. Do not rewrite yet. Flag unclear items with [VERIFY].

PROMPT 2 · TRANSFORM

Using the extracted steps, create a plain-English job aid for [ROLE]. Preserve policy meaning. Include numbered steps and escalation triggers.

PROMPT 3 · REVIEW

Create a manager review checklist for the job aid. Include accuracy, missing steps, escalation clarity, and retention guidance.

PROMPT 4 · FINALIZE

Revise the job aid using the review notes. Keep the output concise and label it as draft pending manager approval.

Why chaining works

- Reduces hallucination risk.
- Reviews each stage.
- Catches errors earlier.
- Creates better evidence.
- Preserves source meaning.
- Avoids giant unreviewable outputs.

Separate instructions from source material

When a prompt contains both instructions and source material, separate them clearly. Use tags like `<instructions>`, `<source>`, `<policy>`, `<notes>`, `<output_format>`, `<constraints>`.

XML SANDWICH TEMPLATE

```
<instructions>
You are a banking procedure translator. Convert the source text into a frontline job aid for
[ROLE]. Preserve meaning. Do not add requirements. Flag uncertainty with [VERIFY].
</instructions>

<constraints>
No customer data. No policy changes. No legal conclusions. Output is a draft pending manager
review.
</constraints>

<source>
[PASTE APPROVED SOURCE TEXT]
</source>

<output_format>
  1. Purpose
  2. What changed
  3. Steps
  4. Exceptions
  5. Escalation triggers
  6. Manager checklist
</output_format>
```

Why this works

The tags separate what the AI should do from the material it should analyze. This matters when working with policies, procedures, loan notes, meeting transcripts, regulatory guidance, campaign briefs, complaint summaries, and vendor documents.

Traceable evidence without overconfident prose

For banking work, you want traceability without asking the model to expose long private reasoning.

AVOID

"Think step by step and show all your reasoning."

BETTER

"Provide a concise rationale for each recommendation. Cite the source section or note used. Flag assumptions with [ASSUMPTION] and uncertainties with [VERIFY]."

RATIONALE PROMPT TEMPLATE

For each recommendation, provide:

1. Recommendation
2. Source used
3. Rationale in two sentences or less
4. Assumptions
5. Items to verify
6. Human reviewer

Why this is safer

It produces reviewable evidence without creating long, messy, overconfident reasoning that a banker may mistake for authority.

Show the AI what good looks like

When you need consistency, show the AI examples.

BASIC PATTERN

Here are two examples of good output. Match the structure, tone, and level of detail.

Example 1:

[INSERT]

Example 2:

[INSERT]

Now produce the same type of output for:
[INSERT NEW INPUT]

Best uses in banking

- Adverse-action language training examples.
- Branch service recovery scripts.
- Campaign copy style.
- Executive report summaries.
- Complaint categorization.
- Procedure job aids.
- Call summary formats.

Guardrail: only use examples that are approved, synthetic, anonymized, or safe for the tool being used.

Five areas for any custom assistant or skill

1 · Identity

WEAK

"You are a helpful assistant."

BETTER

"You are a senior community-bank compliance analyst helping draft internal review checklists from approved policy text."

2 · Behavior rules

Always

- Cite uploaded source sections when available.
- Flag uncertainty with [VERIFY].
- Identify human review owner.
- Label outputs as draft unless approved.

Never

- Invent rates, dates, fees, or policy requirements.
- Use customer identifiers.
- Make legal, credit, compliance, or regulatory determinations.
- Present draft output as final approval.

3 · Task scope

"You help convert approved internal procedure text into frontline job aids. You do not change policy meaning or approve procedures."

4 · Output format

"Always output: Purpose, Steps, Exceptions, Escalation Triggers, Review Checklist, Retention Note."

5 · Knowledge context

"Use the uploaded procedure, brand voice guide, or policy manual as the source. If the answer is not in the source, say [VERIFY]."

When the output is weak, upgrade the prompt one line at a time

1. **Sharpen the role.** "You are a banker." → "You are a branch operations manager writing a job aid for new tellers."
 2. **Add context.** Audience, product, process, goal, and what has already happened.
 3. **Specify format.** Tell the AI exactly how to structure the answer.
 4. **Add constraints.** Tell the AI what not to do.
 5. **Ask for options.** Ask for three options when quality matters.
 6. **Add review.** Tell the AI who reviews the output and what they check.
-

SECTION 15 · PROMPT QUALITY CHECKLIST

Confirm before saving or reusing

- ❑ It starts with a specific role.
- ❑ It gives enough context.
- ❑ It defines the output format.
- ❑ It states what not to do.
- ❑ It uses placeholders instead of customer data.
- ❑ It flags uncertain items with [VERIFY].
- ❑ It identifies the human reviewer.
- ❑ It says whether the output is draft, reviewed, or approved.
- ❑ It includes retention guidance if output will be saved.
- ❑ It does not ask AI to make a final decision.

A prompt becomes a skill when it is reusable, controlled, and reviewable

Prompt

A one-time instruction.

Prompt card

A reusable prompt pattern.

Skill

A governed reusable asset with purpose, allowed inputs, prohibited inputs, guardrails, output format, review owner, retention rule, and version history.

Skill builder template

FIELD	RESPONSE
Skill name	
Role	
Purpose	
Recurring task	
Allowed inputs	
Prohibited inputs	
Guardrails	
Expected output	
Human reviewer	
Retention rule	
Risk tier	Green / Yellow / Red
Version	

Controlled delegation, not clever wording

Do not ask AI to be responsible. Ask AI to be useful. The banker remains responsible for the data entered, the source checked, the output used, the customer impact, the review decision, and the retained record.

Good prompting is not clever wording. Good prompting is controlled delegation.

Source: AiBI Foundations curriculum, Module 3 — Reusable Prompt Cards.