

FOUNDATION GUIDE - V2

Prompt Like a Banker

A foundation kit for writing safe, clear, review-ready AI prompts. Ask clearly. Strip sensitive data. Use approved formats. Mark uncertainty with [VERIFY]. Route risky work to human review.

Before a prompt becomes a workflow, make it safe and reusable.

CORE PROMISE

In five minutes, a banker should know the prompt type, allowed source, output format, data boundary, reviewer, and escalation trigger.

The Safe AI Use Checklist teaches the core habits: strip sensitive data, ask clearly, fact-check, and escalate risky decisions. The Workflow SOP Builder documents repeatable AI-assisted workflows before reuse. This guide sits between them: it teaches staff how to write a safe, review-ready prompt before that prompt becomes a shared team asset.

This is not generic prompt engineering. It is banker prompting: use placeholders, choose the right output format, keep AI in a draft-support role, mark uncertainty with [VERIFY], and route higher-risk work to a qualified reviewer.

The Banker Prompt Formula

ROLE

Who should AI assist?

Name the banker role or review lens.

TASK

What should it produce?

Ask for draft support, not a final decision.

SOURCE

What material may it use?

Use public, internal, redacted, or synthetic source material.

FORMAT

What should the answer look like?

Choose bullets, table, checklist, email, memo, SOP, FAQ, or markdown.

CONSTRAINTS

What must it avoid?

No invented facts, sensitive data, decisions, or unsupported claims.

REVIEW

Who checks it before use?

Name the reviewer, [VERIFY] rule, and escalation trigger.

Choose the work before writing the prompt.

PROMPT TYPE	USE IT FOR	BANK-SAFE CAUTION
Draft	First-pass emails, memos, scripts, huddle notes, and internal updates.	Draft only. A human approves before customer, board, examiner, or retained use.
Summarize	Public guidance, approved procedures, de-identified themes, and training material.	No customer data in public tools. Check facts against the source.
Rewrite	Plain-English versions, tone changes, shorter versions, and job aids.	Preserve meaning. Do not soften compliance language or add requirements.
Extract	Action items, dates, risks, missing facts, open questions, and checklist items.	AI extraction is not verification. Reviewer checks source accuracy.
Compare	Draft versus source, policy versions, product options, or vendor responses.	Ask for differences and uncertainties, not final legal conclusions.
Checklist	Review, handoff, QA, training, and launch readiness workflows.	Reviewer confirms steps match the approved procedure.
Critique	Unsupported claims, missing disclosures, unclear steps, and risk flags.	AI flags possible issues. It does not approve or reject.
Scenario	Fictional training examples and role-play practice.	No real customer, employee, vendor, or case details.
Data-to-narrative	Plain-English summaries of approved aggregate metrics.	Use aggregate or de-identified data unless the tool is approved for the data class.
Solution brief	Turning workflow pain into a problem, user, constraints, and test note.	Keep it internal and move buildable ideas into the Banker Builder Brief Kit.

Risk lanes before export



Green

Safe to draft. Human review still required.



Yellow

Use an approved tool, placeholders, and a named reviewer.



Red

Block public-tool prompting. Route to an approved workflow.

Replace real specifics before the prompt.

Use uppercase bracketed placeholders consistently across the bank's prompt library. The point is not to hide the review obligation. The point is to remove real sensitive facts before asking for draft support.

USE PLACEHOLDER	REPRESENTS	REVIEW NOTE
[CUSTOMER]	Customer or member reference	Use fictional names only in training scenarios.
[ACCOUNT_TYPE]	Checking, savings, CD, HELOC, loan, or card	Do not include account ownership details unless approved.
[ACCOUNT_NUMBER]	Account number or identifier	Never paste real account numbers into public tools.
[DATE]	Relevant date or date range	Use only if timing matters and identity does not.
[AMOUNT]	Balance, payment, fee, or transaction amount	Ranges are safer than exact values for examples.
[TRANSACTION_TYPE]	Deposit, ACH, debit card, wire, check, or ATM event	Use pattern-level language where possible.
[BRANCH]	Branch, team, or location	Remove employee or customer specifics.
[PRODUCT]	Product or service	Confirm terms against approved product material.
[NEXT_STEP]	Human action, handoff, or open item	AI can draft the step; the owner confirms it.
[POLICY_SECTION]	Approved policy or procedure reference	Verify section numbers and wording against source.
[REVIEWER]	Human reviewer role or owner	Name the real reviewer before downstream use.
[VERIFY]	Uncertain, missing, unsupported, or source-dependent item	Resolve before final use.



Placeholder warning

Placeholders reduce exposure, but they do not automatically make sensitive facts safe. If the story still reveals a real customer, investigation, SAR, complaint, loan file, fraud case, or security issue, escalate instead of prompting.

Unsafe versus safer prompt patterns.

UNSAFE PROMPT	SAFER PROMPT
Summarize John Smith's account issue and draft a response about why his debit card transaction was denied.	Using only the redacted facts below, draft a customer-service response template. Do not include account numbers, balances, transaction details, eligibility decisions, fraud conclusions, or promises. Mark missing facts as [VERIFY]. Output as a draft pending banker review.
Write the adverse-action notice and decide the best reason to deny this loan.	Using only [HUMAN_PROVIDED_REASONS], rewrite the draft explanation for clarity. Do not add, infer, rank, approve, deny, or decide. Mark unsupported language as [VERIFY].
Improve this real SAR narrative from the case system and make it more persuasive.	Using [SYNTHETIC_CASE_FACTS] only, create a fictional training example organized by who, what, when, where, why, how, missing facts, and reviewer checklist.
Use these complaint notes to tell me what we should offer the customer.	Using [COMPLAINT_SUMMARY], draft a neutral issue summary and list missing facts. Do not determine fault, offer remediation, or promise outcomes. Route to [REVIEWER].

Copy-ready starter prompt

You are assisting [ROLE] with [TASK].

Use only this approved, redacted, or synthetic source material: [SOURCE_DESCRIPTION].

Do not use or request customer names, account numbers, balances, transaction details, SSNs, DOBs, complaint specifics, SAR/AML case facts, authentication details, examiner-sensitive material, passwords, API keys, access tokens, or production system details.

Return the answer as: [OUTPUT_FORMAT].

Constraints: Do not invent facts. Do not make a decision. Do not add requirements. Preserve source meaning. Mark uncertain or unsupported items as [VERIFY]. Label output as draft pending human review.

Human review: [REVIEWER] confirms source accuracy, data handling, unclear items, final-use approval, and evidence retention.

Escalate if [ESCALATION_TRIGGER].

Require the six checks before copying output downstream.

- ❑ Sensitive data removed or replaced with approved placeholders.
- ❑ Prompt asks for draft support, not a final decision.
- ❑ Uncertain items must be marked [VERIFY].
- ❑ Tool is approved for this data class.
- ❑ Output format is clear.
- ❑ Reviewer and escalation trigger are named.

MOVE TO SOP WHEN REPEATED

If the same prompt becomes a team practice, move it into the Reusable AI Working Brief and then the Workflow SOP Builder. A reusable prompt needs an approved tool, allowed inputs, prohibited inputs, reviewer, approval checkpoint, escalation path, and retention rule.

Source basis: NIST AI Risk Management Framework, OWASP GenAI Top 10, Interagency Guidance on Third-Party Relationships, CFPB Regulation B, CFPB Regulation P, FFIEC BSA/AML Manual risk concepts, and AiBI safe-use/control-plane curriculum.

Adapt before adoption: This guide is a training and implementation aid, not legal advice or an examiner-approved policy. Adapt before adopting it to your institution's approved tools, data classifications, retention schedule, review authority, and escalation procedures.