

TEAM HUDDLE CARD · V1.0

The Banker's Red / Yellow / Green AI Use Card

A 10-second guide for deciding what belongs in AI, what needs review, and what does not belong in a public AI tool. Print the next two pages front/back for a team huddle.

The three lanes

The lanes are not about what AI can do. They are about what your bank will let it do without a second pair of human eyes.

Green - usually safe

- Public or sanitized material: training ideas, agendas, discussion questions, template shells.
- Published sources: FFIEC bulletins, regulator press releases, website copy, approved policies.
- Reusable templates with [CUSTOMER], [ACCOUNT], [DATE], and [AMOUNT] placeholders.

Rule: no sensitive data, no confidential records, no decisions, no customer or examiner impact.

Yellow - approved bank tool + review

- Customer-facing draft language after PII, NPI, and confidential details are removed.
- Internal summaries of non-sensitive policies or approved-source documents.
- Operational checklists, procedure drafts, and first-pass analysis verified by a qualified human.

Rule: use only an approved tool for the data and use case. A qualified human reviews before use.

Red - do not use public AI

- Customer records, examiner correspondence, confidential supervisory information, or security controls.
- Credit decisions, adverse action, fraud, AML, sanctions, regulated reporting, or legal/compliance conclusions.
- Any AI use that approves, denies, ranks, classifies, or decides something requiring documented governance.

Rule: high data risk or decision risk is red even when the audience is internal.

The review rule

If an AI output could affect a customer, an examiner, a policy, a control, or a financial decision, a qualified human must review it before use. No exceptions.

What "qualified human review" means

- The reviewer has the authority to sign off on the work in its non-AI form.
- The reviewer compares the AI output against a known source — your policy, your data, your written procedure — not against memory.
- The reviewer flags uncertain claims with [VERIFY] and does not pass them downstream.

Three questions to ask before you paste

1. **Whose data is in this prompt?** If the answer involves a real customer, escalate or strip.
2. **Which tool is being used?** If the tool is not approved for this data class and use case, classify it red.
3. **What happens if the AI is confidently wrong?** If the answer is "we lose a customer," "we miss a control," or "we file an inaccurate report," classify it red and use an approved internal workflow.

What "approved tool" means

An approved tool is named by the bank for that data class and use case, covered by vendor/data terms, access controls, retention expectations, and a review workflow. Public chat tools are not approved for sensitive or decisioning work unless your bank has explicitly approved that exact path.

When in doubt

Classify higher. Internal-only may reduce audience risk, but data risk and decision risk can still make the use case red.

Next step

Download the companion Safe AI Use Checklist at AIBankingInstitute.com/resources, then run a 10-minute team huddle using one green, one yellow, and one red example from your own department.

Source basis: informed by SR 26-2 revised model risk management guidance; the AIEOG AI Lexicon and Financial Services AI Risk Management Framework released by Treasury, FBIIC, and FSSCC in February 2026; Interagency Third-Party Risk Management guidance; and ECOA / Regulation B expectations for credit-related uses. Starter artifact only; not legal advice or an examiner-approved policy. Adapt before adoption.