



The AI Banking Institute
Turning Bankers into Builders

TEAM HUDDLE CARD · V1.0

Red, Yellow, Green: a one-glance AI use card

A two-page team huddle card. Print it. Tape it next to the coffee machine. Everyone in your bank should be able to classify an AI use case in under ten seconds.

The three lanes

Every use case your team brings to AI fits in one of three lanes. The lanes are not about what AI **can** do — they are about what your bank will let it do without a second pair of human eyes.

Green — usually safe with no sensitive data

- Drafting internal memos that do not name customers.
- Brainstorming training scenarios and discussion questions.
- Summarizing public documents — FFIEC bulletins, regulator press releases, your own published policies.
- Building reusable templates with [CUSTOMER], [ACCOUNT TYPE], [DATE] placeholders.

Rule: if you'd be comfortable putting it on a whiteboard in the break room, it is green.

Yellow — approved tool only, strip sensitive data, human review before use

- Customer-facing draft language — emails, voicemail scripts, branch signage copy.
- Internal summaries of policies or examiner correspondence.
- Operational checklists and procedure drafts.
- First-pass financial or risk analysis where AI is the calculator, not the decision-maker.

Rule: if a customer or examiner could read it, it is yellow. A qualified human signs off before it leaves the building.

Red — do not use public AI tools

- Customer PII, NPI, account records, balances, transaction-level data.
- Credit decisions, adverse action reasoning, or anything in the ECOA / Reg B path.
- Legal, regulatory, or compliance determinations.
- Authentication, fraud workflows, or any record tied to a security control.

Rule: if a regulator would expect a documented model under SR 11-7, it does not belong in a public chat tool.

The review rule

If an AI output could affect a customer, an examiner, a policy, a control, or a financial decision, a qualified human must review it before use. No exceptions.

What "qualified human review" means

- The reviewer has the authority to sign off on the work in its non-AI form.
- The reviewer compares the AI output against a known source — your policy, your data, your written procedure — not against memory.
- The reviewer flags uncertain claims with [VERIFY] and does not pass them downstream.

Three questions to ask before you paste

1. **Whose data is in this prompt?** If the answer involves a real customer, escalate or strip.
2. **Who is the audience for the output?** Internal-only is yellow at worst. Customer or examiner is yellow with mandatory review.
3. **What happens if the AI is confidently wrong?** If the answer is "we lose a customer," "we miss a control," or "we file an inaccurate report," classify it red and use an approved internal workflow.

CITATIONS

This card aligns with the AIEOG AI Lexicon definitions for AI governance, human-in-the-loop, and third-party AI risk (US Treasury / FBIIC / FSSCC, Feb 2026), and with the Federal Reserve guidance on model risk management (SR 11-7).

When in doubt

Classify higher. A yellow use treated like a green use is how banks end up in front of examiners. A red use treated like a yellow use is how banks end up in the press. There is no penalty for being conservative with a tool your team is still learning.