

A 20-minute AI literacy primer for your team

The minimum vocabulary every banker needs before policy, before tools, before anything.

You cannot ask staff to use AI safely if half of them think it is magic and the other half think it is autocomplete. A short, shared vocabulary closes the gap faster than anything else.

This primer is meant to be read aloud at a stand-up or shared as a one-pager in a team meeting. It defines four terms.

The four words every banker should be able to use

Hallucination. When an AI tool produces an answer that sounds right but is fabricated. The model does not know it is wrong. The banker is the only person who can catch it.

Context window. What the AI can "see" right now. Anything outside that window — a document you did not paste, a fact from last quarter — is not part of the answer. If the answer feels strangely confident, ask: did I give it the right context?

Prompt. The instruction you give the tool. A vague prompt produces a vague answer. A precise prompt — role, task, constraints, audience — produces something useful. Prompt quality is your job, not the tool's.

Human in the loop. A banker reviews and accepts every AI output before it leaves the institution. No AI output goes to a customer, a regulator, or a board paper without a human signature. This is not optional.

Three things you can do this week

1. **Read the primer aloud at one team meeting.** Twenty minutes.
2. **Write the four definitions on a wall in the back office.**
3. **Add the four terms to your next monthly all-staff email** with one sentence each.

A starter prompt to test understanding

Ask three colleagues this week: "What is an AI hallucination, in your own words?" If you cannot get three good answers, you have not built shared vocabulary yet — and that is the real work.

Citations

- AIEOG AI Lexicon, US Treasury / FBIIC / FSSCC, February 2026 — official definitions for hallucination, AI governance, HITL
- SR 11-7 Guidance on Model Risk Management, Federal Reserve / OCC