

STARTER TEMPLATE · ADAPT BEFORE ADOPTION

CDFI Grant AI Evidence Checklist

A mission-first, audit-ready way to document where AI assisted your grant, certification, and impact files — so human judgment, fair-lending integrity, and community outcomes stay clearly attributable and defensible.

WHO IT IS FOR

CDFI, MDI, community development, grants, and impact teams

FILL TIME

12 min

WHAT'S INSIDE

- 01 Why an AI evidence file matters for mission lenders
- 02 Build (and reuse) an AI use-case inventory
- 03 Separate AI assistance from human judgment in the record
- 04 Fair-lending and access guardrails (ECOA / Regulation B)
- 05 Map AI evidence to CDFI Fund reporting instruments
- 06 Worked example — AI-assisted annual impact narrative
- 07 Retention, access, and audit-readiness
- 08 What good looks like / common mistakes

SECTION 01

Why an AI evidence file matters for mission lenders

Your CDFI certification, awards, and impact reporting rest on a chain of attestations to the CDFI Fund and, ultimately, to the communities you serve. When AI assists the drafting or analysis behind those attestations, an evidence file proves a human owned the judgment — protecting both your certification and your mission credibility. This is a mission-integrity control, not a tech disclosure: the question a reviewer or examiner asks is "who decided, and on what basis?"

- ^a Scope it to documents that feed obligations: the ACR, the TLR, Performance Progress Reports for applicable awards, impact narratives, and community-development outreach summaries.
- ^a Recommended: keep one lightweight evidence row per AI-assisted deliverable — enough to reconstruct what AI touched and who verified it — rather than logging every keystroke.
- ^a Adopt-verbatim file-purpose statement: "This evidence file records where AI tools assisted in preparing materials submitted to or relied upon for the CDFI Fund and for community-impact reporting. AI assistance does not replace the certifying official's independent review; all figures, claims, and certifications reflect human verification."
- ^a Not legal advice — confirm specifics against your Award/Assistance/Allocation Agreement and your award reporting instructions.

SECTION 02

Build (and reuse) an AI use-case inventory

Before logging individual documents, list the AI uses you actually permit in mission work. A short inventory turns ad hoc tool use into governed, repeatable practice and gives you a stable vocabulary for every downstream evidence row. The AIEOG Lexicon frames AI governance as the policies, roles, and oversight that direct how AI is adopted and monitored.

- ^a Recommended default prohibitions for mission lenders: no AI-generated eligibility or adverse-action reasoning, no synthetic statistics in impact claims, and no AI text inserted into a certification without a named reviewer's sign-off.
1. List each approved AI use case in plain language (e.g., "draft impact-narrative first version," "summarize small-business outreach notes," "check ACR narrative for internal consistency").
 2. For each, record the tool/model, the data category it may touch, and the explicit prohibition line (e.g., "no borrower PII," "no protected-class inferences," "no auto-generated certification language submitted without review").
 3. Assign a human owner accountable for each use case and the verification step required before the output is used.
 4. Note where outputs land (ACR narrative, PPR, impact report, outreach summary) so the inventory cross-references your evidence rows.
 5. Review the inventory at least annually — align the cycle with your CDFI Fund reporting calendar so it is current when attestations are due.

SECTION 03

Separate AI assistance from human judgment in the record

The single most valuable thing your file does is draw a clean line between what AI produced and what a person decided — the same principle CFPB enforces in lending (an institution cannot hide behind a model), applied to documentation.

- ^a Default rule: every AI-assisted deliverable carries a one-line attribution note naming the AI role, the human reviewer, and the verification performed.
- ^a Adopt-verbatim reporting note (working files or cover memo): "Portions of this document were drafted with AI assistance. All data, eligibility determinations, impact claims, and certifications were independently reviewed and verified by [name/title] on [date]. The institution takes full responsibility for the accuracy of the final content."
- ^a Distinguish three AI roles so reviewers know the stakes: drafting (low — language only), summarizing (medium — verify against sources), and analysis/recommendation (high — requires documented human re-derivation).
- ^a Never let AI generate the certification or attestation language itself; the certifying official's words and accountability must be human-authored.

SECTION 04

Fair-lending and access guardrails (ECOA / Regulation B)

Mission lending lives or dies on equitable access, and AI introduces concrete fair-lending exposure. Per CFPB guidance, ECOA and Regulation B do not permit creditors to use technology so complex they cannot provide specific and accurate reasons for an adverse action — complexity is not an excuse. Your file should show AI never became a black box between an applicant and a specific, accurate reason.

- ^a Recommended: AI is permitted for drafting and summarizing, but credit decisions, eligibility calls, and adverse-action reasons must be human-determined and independently documented.
- ^a Do not rely on generic checkbox reasons when AI or complex models are involved; reasons must accurately describe the factors actually considered.
- ^a Log a fair-lending check on any AI-assisted analytical output: confirm no protected-class proxy entered the inputs and that conclusions are reproducible by a human without the tool.
- ^a Adopt-verbatim guardrail note: "No AI output was used to determine applicant eligibility, pricing, or adverse-action reasons. AI assistance was limited to language drafting and summarization of human-verified content; all decisioning rationale is human-authored and specific to the individual circumstances."

SECTION 05

Map AI evidence to CDFI Fund reporting instruments

Your obligations flow through named instruments and your individual agreement; AI-assisted preparation of any of them belongs in the evidence file. The point is traceability from the submitted attestation back to the human who verified it.

- ^a ACR (Annual Certification and Data Collection Report): if AI helped draft narrative responses, log the row and confirm the certifying official reviewed every certification statement.
- ^a TLR (Transaction Level Report): AI must not generate or alter transaction-level data — if used at all (e.g., formatting QA), record it explicitly and verify against source systems.
- ^a Performance Progress Report and other award reports: for active recipients, log AI assistance on aggregate performance narratives and reconcile figures to systems of record.
- ^a Material Events: where AI assisted in drafting a notification, ensure timeliness (commonly within 30 days, or as your agreement specifies) and human verification of the underlying facts.
- ^a Recommended: for anything submitted to the CDFI Fund, the evidence row must name a human certifier; when an instrument or deadline is uncertain, defer to your grant agreement and reporting instructions and the CCME Help Desk rather than assuming.

SECTION 06

Worked example — AI-assisted annual impact narrative

Here is the practice in miniature. Use this filled row as your template.

- ^a Scenario: the Director of Impact uses an approved AI tool to turn de-identified, aggregate small-business lending data and field notes into a first-draft impact narrative for the certified Investment Area, plus a one-paragraph outreach summary from staff notes.
- ^a Verification performed: the Director reconciled every figure to the loan-origination system, removed two AI phrasings that overstated job-creation outcomes, confirmed no borrower-level or protected-class data was used, and the certifying official independently approved the final narrative.
- ^a Sample filled evidence row — Document: FY2025 Annual Impact Narrative (supports ACR) | AI role: first-draft drafting + summarization | Tool: [approved internal AI assistant] | Data: de-identified aggregate program data; staff field notes | AI prohibited from: borrower PII, certification language, impact statistics | Reviewer: Director of Impact | Verification: figure-to-source reconciliation; overstatement removed; fair-lending data check passed | Certifying sign-off: CEO, 2026-03-04 | Retention: 5 years from submission (per agreement).

SECTION 07

Retention, access, and audit-readiness

Evidence is only useful if it survives staff turnover and can be produced on request. Tie retention to your reporting cycle and keep the file simple enough that anyone can reconstruct what happened.

- ^a Recommended retention: keep AI evidence rows and supporting drafts for the longer of (a) the period required by your Award/Assistance/Allocation Agreement or (b) your standard records-retention schedule — commonly several years past the related submission; confirm the exact term in your agreement.
- ^a Store evidence rows alongside the submitted instrument (ACR/TLR/PPR/impact report) so a reviewer can move from attestation to verification in one step.
- ^a Control access: restrict the file to staff with a need to know, and never store borrower PII or protected-class data in it.
- ^a Adopt-verbatim retention note: "AI evidence records for this reporting period are retained for [X] years consistent with our records-retention policy and our agreement with the CDFI Fund, and are available to authorized reviewers and examiners upon request."

SECTION 08

What good looks like / common mistakes

A strong AI evidence file reads like a mission–integrity ledger: clear human ownership, clean separation of AI from judgment, and no fair–lending or access blind spots. Most failures come from over– or under–documenting in the wrong places.

- a** What good looks like: every AI–assisted submission has a named human certifier, AI's role is described in one honest line, figures are reconciled to source, and fair–lending/access checks are recorded.
- a** Common mistake (transparency): omitting AI involvement entirely, or burying it so a reviewer cannot tell what AI touched.
- a** Common mistake (fair lending): letting AI generate or influence eligibility or adverse–action reasoning, then being unable to give a specific, accurate, human–verifiable reason.
- a** Common mistake (access/accuracy): letting AI inflate impact or reach figures without reconciling to source data.
- a** Common mistake (over–documentation): storing borrower PII or protected–class data in the evidence file itself — keep the ledger about the process, not the people.

SOURCE BASIS

- CDFI Fund (US Treasury) — CDFI Certification reporting: Annual Certification and Data Collection Report (ACR) and Transaction Level Report (TLR)
- CDFI Fund — award compliance and reporting obligations (Assistance/Allocation Agreement, Performance Progress Report, Material Events; CCME Help Desk)
- CFPB guidance on adverse–action notices when using AI/complex credit models (ECOA / Regulation B; Circular 2022–03)
- 12 CFR Part 1805 — Community Development Financial Institutions Program
- AIEOG AI Lexicon — AI governance, AI use case inventory (US Treasury / FBIIC / FSSCC, Feb 2026)
- Your grant agreement, award conditions, and reporting instructions; institution records–retention policy

ADAPT BEFORE ADOPTION

This is a starter template, not legal advice or an examiner–approved policy. Replace every [bracketed] placeholder, align the language with your institution's policies, and route the final document through compliance, risk, and legal review.