

STARTER TEMPLATE · ADAPT BEFORE ADOPTION

Go-to-Market Plan for an AI Initiative

A run-it-this-week go-to-market plan for launching an AI-enabled capability at a community bank or credit union — internal rollout plus a compliant member-facing announcement — with a worked example threaded through every section.

WHO IT IS FOR

Marketing, retail, and product leaders launching an AI capability internally or to members

FILL TIME

12 min

WHAT'S INSIDE

- 01 Define the launch and the audience
- 02 The promise (positioning & value proposition)
- 03 The proof (evidence, pilot & substantiation)
- 04 Channels & messaging (internal + external)
- 05 Compliance review & guardrails
- 06 Timeline & rollout sequence
- 07 Metrics & success criteria
- 08 What good looks like / common mistakes

SECTION 01

Define the launch and the audience

A GTM plan fails when "AI" is the product instead of the outcome. Name one concrete capability, the job it does, and exactly who you're rolling it out to internally and announcing to externally — narrow scope is what makes a launch compliant, measurable, and finishable.

- ^a Recommended default: pick an internal-efficiency or human-in-the-loop use case first, not a member-facing autonomous one — the 2026 differentiator is who operationalizes AI with discipline, not who ships the flashiest feature.
- ^a Split the audience into two tracks: the internal audience who must adopt it (for the worked example, ~25 contact-center agents, 3 team leads, 1 compliance reviewer) and the external audience who will hear about it (members who contact support).
- ^a Decision rule: announce to members only when the AI materially changes their experience or when silence would be misleading; a back-office drafting aid a human approves usually needs internal change management, not a press release.
- ^a Adopt-verbatim internal framing line: "This tool drafts; you decide. Nothing reaches a member until a person on our team reads it and approves it."

SECTION 02

The promise (positioning & value proposition)

Your promise is the one sentence that justifies the launch to the people who must adopt it and the members who experience it. It should describe a member or staff outcome, never the technology, and it must be something you can actually deliver and substantiate.

- ^a Write it as "[audience] gets [outcome] because [what we changed]" — outcome first, mechanism second; the word "AI" is optional and often better left out of member-facing copy.
- ^a Internal promise for the worked example: "Agents resolve member questions faster and with more consistent, on-brand answers, because AI drafts a first reply they can edit in seconds instead of writing from scratch."
- ^a Member-facing angle (when warranted): lead with responsiveness and personal service — the attributes research identifies as why consumers switch to community institutions — anchored to your existing trust advantage, not novelty.
- ^a Hard constraint: every adjective must be provable. "Faster" needs a baseline and a measured delta; "more accurate" needs QA data. If you can't measure it, don't claim it.
- ^a Adopt-verbatim member-facing line (use only if you announce): "We've added new tools that help our team answer your questions more quickly — and a real person on our team still reviews every response before you get it."

SECTION 03

The proof (evidence, pilot & substantiation)

Proof is both a marketing asset and a compliance shield. Under UDAAP, any performance claim must be substantiated before you make it — so the pilot that proves the value is the same data that protects the claim. No claim ships ahead of its evidence.

- ^a Recommended default: do not publish a numeric claim ("answers 30% faster") unless the pilot produced that exact number under documented conditions; if directional, make the qualitative point with no figure.
- ^a Adopt-verbatim internal talking point: "Before we tell a single member anything, we have to show the numbers behind it. The pilot is how we earn the right to make the claim."
- 1. Run a time-boxed pilot (2–4 weeks) with a subset of the internal audience — for the worked example, 5 agents using AI-drafted replies while the rest serve as a control group.
- 2. Capture a clean before/after baseline on the metrics you intend to claim: handle time, first-contact resolution, QA accuracy/error rate, and satisfaction.
- 3. Have compliance/QA review a sample of AI-drafted-then-approved replies for accuracy, tone, and prohibited claims (e.g., "instant approval," "no fees," "best rate") — the failure mode the CFPB chatbot spotlight flagged.
- 4. Lock the substantiation file: the specific numbers, date range, sample size, and who approved them — this authorizes any external claim and is what you produce if examined.

SECTION 04

Channels & messaging (internal + external)

Channels split by audience: internal channels drive adoption and confidence; external channels set member expectations. Match the message to the channel and never let an external claim outrun your substantiation file.

- ^a Internal channels for the worked example: a hands-on training session, a one-page "draft, edit, approve" quick reference, a pinned chat message, and a standing huddle item for the first month.
- ^a External channels (only if announcing): in-app/online-banking message, a short FAQ or help-center article, branch talking points, and a plain-language explanation for any member who asks.
- ^a Recommended disclosure default: be transparent and plain — consumer comfort with AI rises when a specific benefit is explained and a human-oversight role is named; vague "AI-powered" labeling erodes the trust community institutions depend on.
- ^a Adopt-verbatim branch/phone talking point: "Yes — our team uses tools that help us draft answers faster, but a person here always reviews and approves what you receive. Your information stays protected and you can always reach a real person."
- ^a Hard rule: do not imply members are talking to a human when they're talking to a bot, and do not imply a bot can do something it can't — both are direct UDAAP risks.

SECTION 05

Compliance review & guardrails

Compliance is a gate in the plan, not a sign-off at the end. Build the review into the timeline so compliance shapes claims and disclosures before launch, with extra scrutiny on anything touching credit. (Operational guidance, not legal advice — route final decisions through your compliance counsel.)

- ^a Map the use case to its risk tier: back-office drafting with human approval (lower risk) vs. anything that influences a credit decision, eligibility, pricing, or adverse action (high risk, additional rules apply).
- ^a Scope guardrail for the worked example: AI-drafted replies must not state credit decisions, approval odds, rates, or eligibility — lending questions route to the proper process rather than letting a drafted reply make a representation.
- ^a Reg B / ECOA guardrail: CFPB Circulars 2022-03 / 2023-03 require specific, accurate reasons for adverse action regardless of technology; a creditor may not hide behind black-box AI. Keep AI out of adverse-action messaging unless you can meet this.
- ^a Adopt-verbatim agent escalation rule: "If a drafted reply mentions a rate, an approval, eligibility, or a denial, do not send it — escalate it. AI never delivers a credit decision."
- ^a Recommended: require compliance sign-off on three artifacts before launch — the member-facing announcement copy, the FAQ/disclosure language, and the substantiation file behind any claim.

SECTION 06

Timeline & rollout sequence

A realistic first launch runs roughly 6-10 weeks. Sequence it so compliance and substantiation precede any external word, and so internal adoption is solid before members are involved.

- ^a Gate rule: no external announcement ships until the substantiation file is signed and internal adoption is confirmed — an unevenly adopted tool plus a public promise invites complaints.
1. Weeks 1-2 — Define & align: lock scope, audiences, the promise, success metrics, and the prohibited-claims list; open the compliance workstream.
 2. Weeks 3-4 — Pilot: run it, capture baseline vs. results, assemble the substantiation file.
 3. Week 5 — Compliance review & decision: review pilot output and draft copy; decide whether a member announcement is warranted; sign off on the three artifacts.
 4. Weeks 6-7 — Internal rollout: train the full internal audience, publish the quick reference, require the draft/edit/approve workflow.
 5. Week 8 — External announcement (if warranted): publish in-app message and FAQ, brief branch staff, turn on monitoring.
 6. Weeks 9-10 — Monitor & adjust: track metrics and complaints weekly; fix or pull anything that misses the bar before scaling.

SECTION 07

Metrics & success criteria

Define success before launch or you'll declare victory by anecdote. Pick a small set of metrics tied to the promise, set baselines, and name the threshold that would make you pause, fix, or roll back.

- ^a For the worked example, track four: average handle time, first-contact resolution, QA accuracy/error rate on sent replies, and member CSAT — plus a guardrail metric: AI-related complaints or escalations.
- ^a Set the success threshold in advance: e.g., "handle time down with no drop in QA accuracy and no rise in complaints"; a speed gain that increases errors is a failure, not a win.
- ^a Define a kill/rollback trigger: e.g., QA error rate rises above the pre-launch baseline, or member complaints referencing inaccurate/"robotic" responses exceed a set count in a week — pull or pause and remediate.
- ^a Adopt-verbatim leadership update line: "We measured it before we scaled it: [metric] moved from [baseline] to [result] with no increase in errors or complaints — here's the data."

SECTION 08

What good looks like / common mistakes

The difference between a launch that compounds and one that creates risk is discipline: scoped capability, substantiated claims, human oversight, transparent disclosure, and a metric that can fail.

- ^a What good looks like: one narrow capability with a human in the loop; a promise stated as an outcome; every claim backed by a dated substantiation file; plain-language disclosure naming the benefit and the human-review role; compliance signed off before any external word; a defined success threshold with a rollback trigger.
- ^a Common mistake (overpromising): claiming "instant," "guaranteed," "best rate," or a percentage you didn't measure — the CFPB chatbot spotlight flags AI generating unsubstantiated specifics, which map straight to deceptive-acts claims.
- ^a Common mistake (UDAAP): marketing that doesn't match the fine print, implying a member is talking to a human when it's a bot, or omitting AI's role where a member would have needed to know it.
- ^a Common mistake (Reg B/ECOA): letting AI touch credit messaging or adverse-action reasons without being able to give specific, accurate reasons; "the model decided" is not a permissible reason.
- ^a Common mistake (no metric): launching without a baseline, threshold, or kill switch, so you can't tell whether it worked or defend the claim later.
- ^a Common mistake (leading with the tech): marketing "AI" instead of the outcome, trading away the trust advantage that is exactly why members choose community institutions.

SOURCE BASIS

- Jack Henry Strategy Benchmark (AI named a top planned technology investment among community banks and credit unions)
- Cornerstone Advisors, "What's Going On in Banking" (AI adoption; no AI strategy without a credible data strategy)
- Apiture / Harris Poll consumer banking studies (personalized, responsive service as a switching driver to community institutions)
- CFPB Issue Spotlight, "Chatbots in Consumer Finance"
- CFPB Circulars 2022-03 / 2023-03 — adverse-action notices for AI/complex credit models (ECOA / Regulation B)
- UDAAP examination guidance (CFPB / NCUA / FDIC / OCC)

ADAPT BEFORE ADOPTION

This is a starter template, not legal advice or an examiner-approved policy. Replace every [bracketed] placeholder, align the language with your institution's policies, and route the final document through compliance, risk, and legal review.